



Oles Shyshkotsov, 4A Games

Sergei Karmalsky, 4A Games

Benjamin Archard, 4A Games

Dmitry Zhdan, NVIDIA



# EXPLORING RAYTRACED FUTURE IN METRO EXODUS

“My old dream was to see global illumination in an interactive application, which doesn't depend on any precomputation and works with 100% dynamic lighting conditions and a similarly dynamic environment.”

Oles Shyshkovtsov, 4A Games (GPU Gems 2, 2005)



# AGENDA

1. Introduction
2. Implementation
3. Denosing
4. Artist point of view



# INTRODUCTION

The Quest for the Holy Grail



MOTIVATION  
RTX OFF





MOTIVATION  
RTX ON



# THE QUEST BEGINS

Know what you want to achieve

Everything?

Ok, fine. Global Illumination!

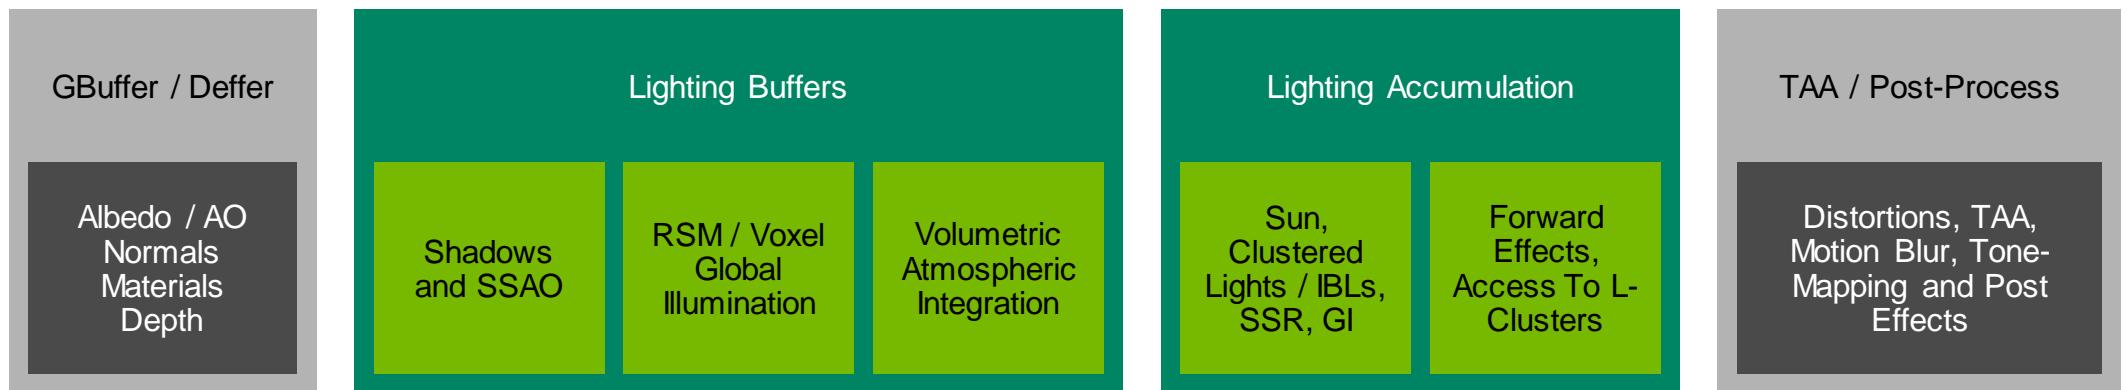
Ok, fine. A hybrid, indirect first-bounce, diffuse, Global Illumination and Deferred rendering pipeline.

Right. Off you go.



# A CUTTING EDGE ENGINE

## Legacy version



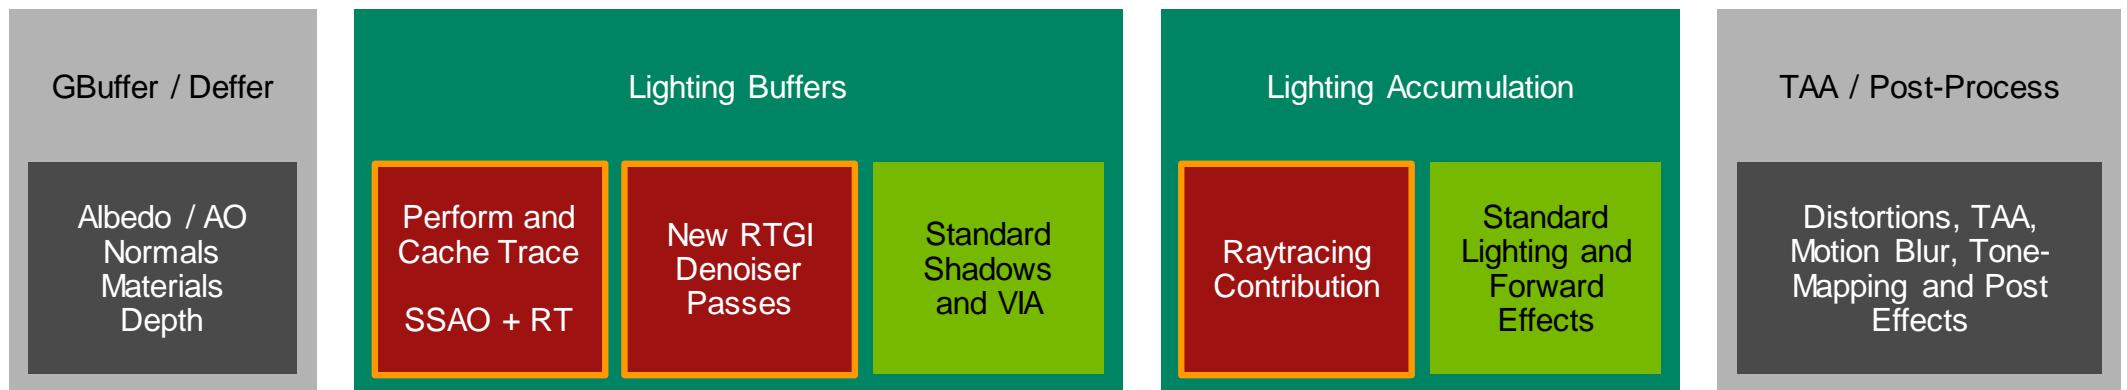
A standard deferred renderer

Calculate G-Buffer and lighting buffers, accumulate light and effects, TAA, post-process

Relies heavily on stochastic methods followed by TAA to reduce noise

# A CUTTING EDGE ENGINE

## RTX ON



The raytraced elements merge nicely with standard deferred renderer

Added buffer to cache raytrace data for use in RTAO and RTGI passes

Stochastic ray generation: Image is figuratively shouting at us, so a good denoiser is critical

# STOCHASTIC METHOD

Why do it 1000 times when once will do?

Monte Carlo Integration: the approximation of large data sets with few samples

Rather than adding up every single ray/photon, pick a few as “representative”

You may lose some specific details, but will get the big picture

Desirable for a GI solution

Results get to the point of diminishing returns at a tiny fraction of the total set

This doesn't just apply to raytracing

Shadows, volumetrics, reflections, hair

Apply to any suitably complex effect

# STOCHASTIC METHOD

## The importance of samples

Beware of noise and aliasing, both are issues, aliasing is worse

You are going to produce a noisy data set, and you will run denoisers

Jitter, Importance sampling and Probability Density Function (PDF) provide leverage over sample distribution

Output data is buffered for analysis, filtering and use later on

Produces good general purpose (input agnostic) data on scene illumination

If your image is still noisy at the end of the frame (it will be) add TAA

# STOCHASTIC METHOD

## Know your noise

Noise breaks up patterns when sampling below input frequency

Must be repeatable, it is used later for re-construction of the hit location from stored distance value

Temporally and spatially uniform to avoid “clumping” and “swimming”

Sample small blue noise texture across the screen, oscillate across frames

```
// Sample noise from screen position and frame index
float2 uv          = t_blue_noise_64.Load(uint4(
    pixID.xy+32*(frmID&1))&63, frmID&63, 0).xy;

// Generate ray on hemisphere
float3 vRay        = HemisphereSample(uv);

// Transform to local space of the surface using
// surface normal
float3 T, B;      BasisFromDirection(N, T, B);
return            normalize(FromLocal(ray, T, B, N));
```

# RTAO

## Why?

AO is a poor-man replacement for GI. We are doing real GI already so why bother?

We are running hybrid pipeline, which is smoothly blended into “old” pipeline

250m transition from foreground RTGI to “regular” pipeline

Regular pipeline expects AO available at different stages

All image-based lighting (light-probes) are directly multiplied by AO

Some “fake” lights use AO as their shadow approximation. Shame on us :)

Even sun shadow-map blends into AO at some distance

Searching for usage in shader-code finds 79 places...

Also, it’s cheap and helps guide the denoiser! :)

# SSAO TO RTAO

## Reuse and Improve



SSAO  
Captures nearby details



RTAO  
Recognise enclosed space



SSAO  
Misses interior occlusion



RTAO  
Progressively darkens

# VOXEL-BASED GI TO RTGI

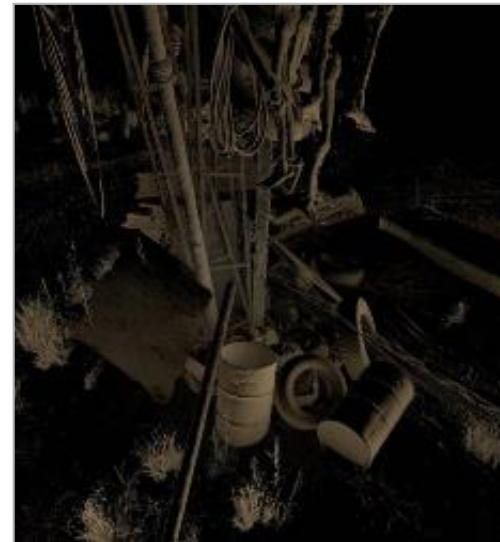
## Night vs Day



**Voxel GI**  
Broad directional light,  
insufficient detail for  
shadows



**RTGI**  
Light bounce and contact  
shadows from nearby  
objects



**Voxel GI**  
No sense of depth



**RTGI**  
Gradual self-occlusion on  
object interiors







RTX OFF



RTAO PASS



RTGI PASS



RTX ON



RTX OFF

RTAO PASS



RTGI PASS

RTX ON

# PUTTING IT ALL TOGETHER

Ok. Now it works. Just...

RTX fitted in well with 4A engine

The game was balanced for the traditional pipeline, but RTX walked in made it its own

We want more “rays”: We generate as few as possible for performance, but we can always find as use for more them

Lots of options for the future...





# IMPLEMENTATION

It WILL just work, if you work at it

# IMPLEMENTATION (1)

Remove unnecessary pipeline stages

RSM rendering (replaced with cheaper depth-only shadow-map rendering)

Geometric ESM-AO (approximation of 16 rays)

SH-voxel-grid computation/gather

SH-voxel-grid temporal blending

SH-voxel-grid screen-space resolve

# IMPLEMENTATION (2)

## Modify some pipeline stages

SSAO-pass now computes accumulation weights and accumulates raytraced AO

Velocity, depth disocclusion, etc.

Weights used for both AO accumulation and GI

AO-filter pass

Before:

SSAO filtering, geometric ESM-AO sampling

Blending with terrain AO, precomputed AO maps, per-vertex AO

Now:

Denoising and RTAO accumulation

# IMPLEMENTATION (3)

## Add new pipeline stages

Raytracing  + screen-space pre-tracing

Geometry skinning and animation

Albedo updates/management

BLAS updates

TLAS rebuilds

Deferred shading of hit-positions

Denoising & accumulation

# RT MODULE

Separate mini-engine from the rest of the pipeline

Handles skinning and geometric animation

Handles all BLAS updates/TLAS rebuilds

Separate instance-culling (expanded frustum, contribution)

Instance transforms, logical/game visibility

Separate memory manager

Separate command lists

Just 3 .cpp files, ~1500 lines DXR API, ~1100 lines logic, ~200 lines “glue”

# BVH MANAGEMENT

## Recipe

**BLAS = update only** for skinned/animated instances; **TLAS = rebuild only** from scratch

TLAS quality and compactness is extremely important

TLAS selects those which are inside expanded frustum (+logical visibility, + contribution culling)

Usually we have more than 100k potentially active instances; less than 5k will survive the culling

Relatively fast, but each update/rebuild is multi-pass, under utilizes GPU

Hide with async-compute!

We hide it with pre-trace CS and SSR CS

Alternatively run it from compute queue parallel to the gbuffer rendering

We have both modes implemented, statistically insignificant perf difference

# BLAS UPDATE THROTTLING

## Skinned and animated vertices processing

Every entity update increases priority of RT-instances

Visible = higher priority, small and/or distant = lower priority

Sort instances based on accumulated (across frames) priority

Select a few (16 in our case) with highest priority

Select a few (4 in our case) randomly from the remaining set with non-zero priorities

High priority objects should not block other stuff updating!

Shrinks queue to "balanced" state in a matter of seconds

"Balanced" state is just 5k-6k instances "outdated" :) out of 20k+

Additionally limit the vertex count as well

Necessary to avoid rare "spikes" in processing

# METRO IS EXTREMELY GEOMETRY HEAVY

20,000,000 polygons is a lot to render just for the sun shadow

Depth impostor cache / Simplified IB (separate position-only VB if shader allows)

Reuse those simplified "shadow" meshes for RT!

Result: BLAS meshes are about 4x smaller than the “real” geometry

There are scenes where it translates into 30% perf gain in raytracing

All vertex animation and skinning become cheaper

Memory usage: ~1GB instead of ~4GB

Zero or close to it difference in quality!

# RAYTRACED GI

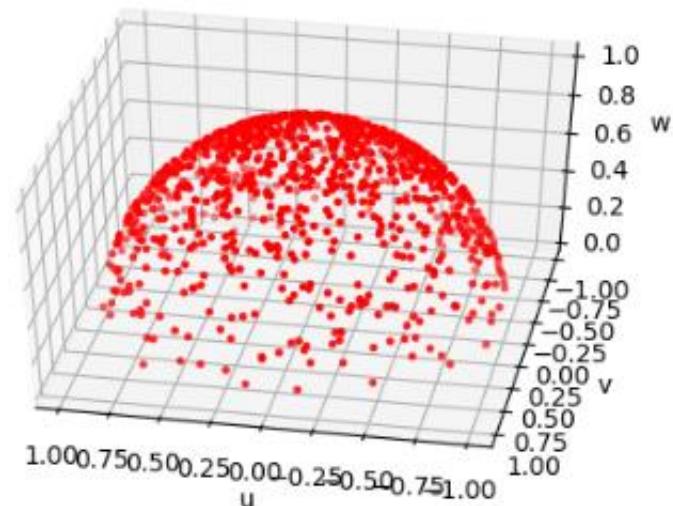
## Basic idea

Shoot rays at every pixel in all directions (ok, according to BRDF lobe)

Gather lighting at the contact point; multiplied by albedo of that point

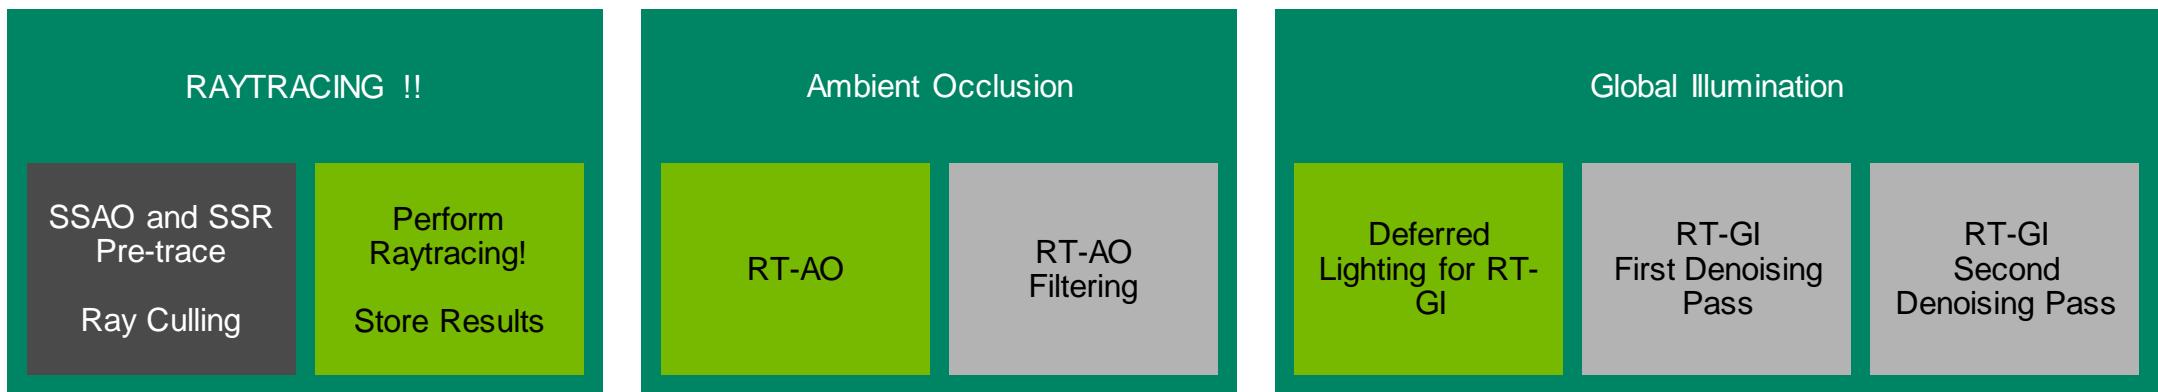
Accumulate that!

Hit distance gives us "free" RTAO



# PIPELINE STAGES

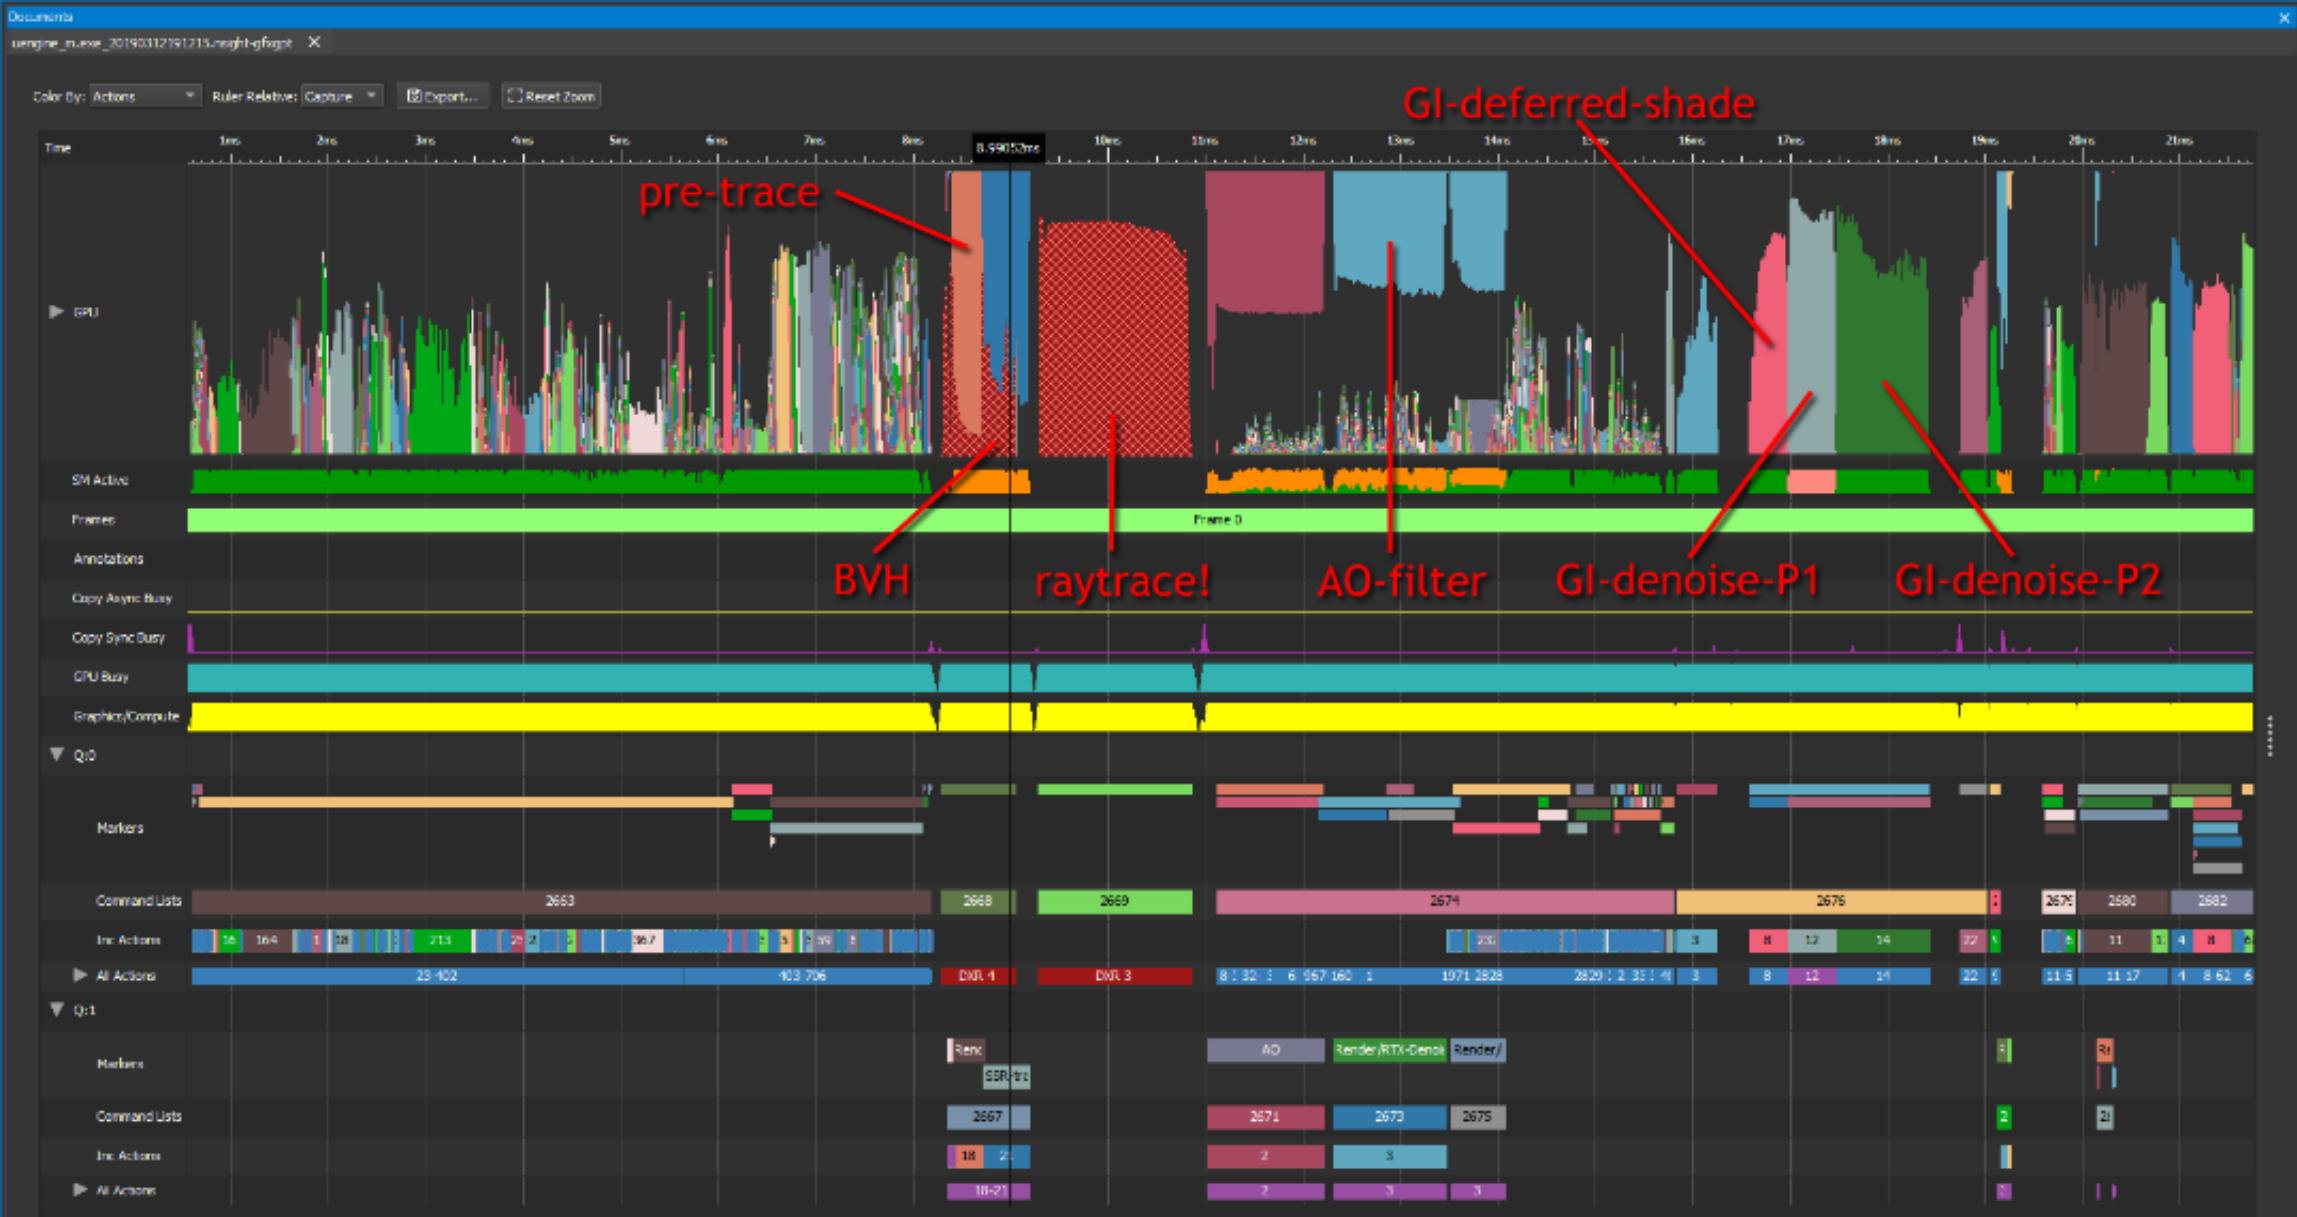
## Raytracing Specific GPU Pipeline



Screen-space pre-trace + all actual raytracing

Ambient Occlusion + Filtering

Global Illumination + Two pass denoiser



# PRE-TRACE

## Ray tracing in screen space

Exactly the same ray-generation as the real raytrace

Ray-march against depth buffer

Runs as async-compute, parallel to BVH updates/rebuilds

Fixes missing "alpha-tested" geometry in most cases

We aggressively filter it out whenever we can

Almost constant distance in screen-space (cache-friendly)

Outputs into UAV hit-distance and albedo (from g-buffer)

# RAYTRACING

## Real rays!

Only spawn the real ray if pre-trace failed to find intersection

- Leads to a small perf-boost

Ray-marches terrain's heightmap inside the "raygen" shader

- Limit ray distance if intersection is found

- Almost free here (if done carefully) due to GPU latency hiding

Extremely simple pipeline config

- Only [shader("closesthit")] is necessary for us to get hit results

- Payload is a single UINT

Outputs to the same UAV, distance + albedo (packed into a single UINT)

- Needs to be careful with precision and tolerances

- Floating point precision hit us several times

# DEFERRED LIGHTING

## Hit-positions processing

Run exactly the same ray generation as in main trace

Reconstruct hit position (or indication of "miss") and albedo

**MISS** = sample skybox

**HIT** = compute lighting

Encode information, more on that later

Accumulate with history

# DEFERRED LIGHTING

Why only the sun/moon and sky?

Tech stabilized quite late in the development cycle (late Q4/2018)

Content was mostly done and locked in at the time

Implemented 1st bounce contribution from all lights, out of curiosity

- Lighting already computed in a deferred way? use it

- In frustum, but occluded? Use precomputed lighting from atmosphere

- Out of frustum - run real computation

Extremely cheap (~0.2ms on an RTX 2080ti), could be a big perf-boost if we managed to remove AO/IBL, but...

- It conflicts with hand-crafted lighting and visuals :(

- It breaks the game, especially the stealth mechanic

Simply put: we were out of time to fix current content across the huge game

# COLOR TRANSPORT

## Where to get albedo for hit results?

Color bleeding is mostly visible on close to contact surfaces

Usually those are found by initial screen space pre-trace

Just sample albedo from gbuffer

Integration across the whole hemisphere is a low-pass filter in essence

It is a good idea to pre-filter signal to lower denoiser's input noise level

We do that pre-filtering extremely aggressively - we store average albedo per-instance :)

Low input noise and extremely fast :)

# COLOR TRANSPORT

Where to get albedo for hit results?

G-buffer (the pre-trace samples this)



Per-instance albedo (raytracing samples this)



# COLOR TRANSPORT

## A few problems

Usually average albedo color pre-calculated per-texture suffices

What to do with metals? Theirs albedo is essentially zero...

Solution: Albedo \* (1 - F0) + F0

What if complex shading changes visible albedo?

Or maybe it is texture-atlas and average doesn't make sense?

Solution: pre-render that exact combination of mesh-shader-textures-params!

Then average visible albedo from 6 directions

Store into sparse database/hash table

Still allow artists to “override” it

Database shipped in the first “hotfix”



# Color bleeding - RTX ON



# IRRADIANCE STORAGE & ENCODING

## Directional color space

Decompose HDR-RGB into Y and CoCg

Encode Y as L1 spherical harmonics (world space), leave CoCg as scalars

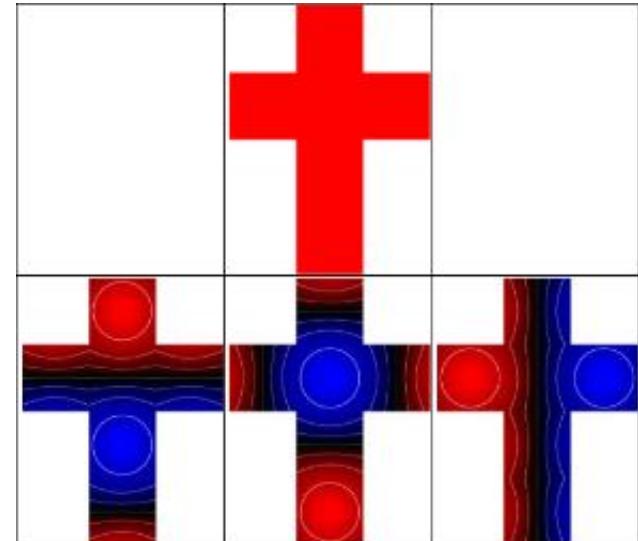
Human eye more sensitive to intensity, not color

4xFP16 for Y

2xFP16 for CoCg

96 bits per pixel in total

All the accumulation and denoising happens in this space



# WHY NOT JUST COLOR?

would R11G11B10 be enough?

Denoisers could go really wide under certain conditions

- Loss of normal-map details

- Loss of "contact" details and general blurriness

- Loss of denoising quality if we weight heavily against normals of samples, less information could be "reused"

96 bits? Why not less?

- Tried to reduce it down to 64 bits - failed

- Mostly because of "recurrent" nature of denoisers which could be extremely aggressive on temporal accumulation and thus precision

- In case of LDR, Y would be in range of [0..1] and CoCg in [-1..1], in our case it is actually in [0..HDR] and [-HDR..+HDR]

# SPECULAR!

## Important for PBR materials consistency

This encoding is actually a low order approximation of cubemap

But at each individual pixel!

This allows us to reconstruct indirect specular!

Crucial for metals where albedo is zero or close to it



( Illustration from paper "Stupid Spherical Harmonics (SH) Tricks" by Peter-Pike Sloan )

# DECODING IRRADIANCE

## Details

Resolve SH as usual against pixel's BRDF to get diffuse

Extract dominant direction out of SH

Compute SH degradation into non-directional/ambient SH

If SH is non-directional - it means incoming light is uniform over hemisphere

And if it is uniform - that's the same as if material is "rough" -> recompute new roughness

Run regular GGX with (extracted\_direction, recomputed\_roughness)

# SPECULAR GI OFF

Booooooooo





**SPECULAR GI ON!**

Yay \(\bullet\)\(\bullet\)/



# THE POWER OF PIPELINE

## Details

The BRDF importance sampling doesn't care what to integrate at all, it is "unbiased" in that sense

Be it 1st, 2nd or 3rd bounce indirect lighting or "direct" lighting or whatever

What if we put something emissive in the scene?

DEMO TIME!



# POLYGONAL LIGHTS

## Details

Yes, that's arbitrary shaped and textured polygonal lights

I saw a lot of research on that...

But nobody does shadows,  
right? 😊

It is free!



# WRAPPING THINGS UP

## “Holy Grail” cracked!

Game-scale realtime 1st bounce indirect lighting from any analytic light

Not limited to 1st bounce at all, but... Xms trace Yms light per bounce

Even 2nd bounce gives diminishing returns compared to cost

Direct lighting and shadowing from arbitrary shaped polygonal area lights

Or sky, or whatever... Artistic freedom...

Computes both diffuse BRDF (Disney) and specular BRDF (GGX)

Everything is fully dynamic, both the geometry and lighting (no precomputation!)

In fact 4A-Engine doesn't really have a concept of something static (prebaked)

Massive scenes

~150 000 000 triangles on a typical Metro level in TLAS before culling



# DENOISING

Trapping the beast in 15 mins

# DENOISING

## What is it?

Denoising (or noise reduction) is the process of removing noise from a signal

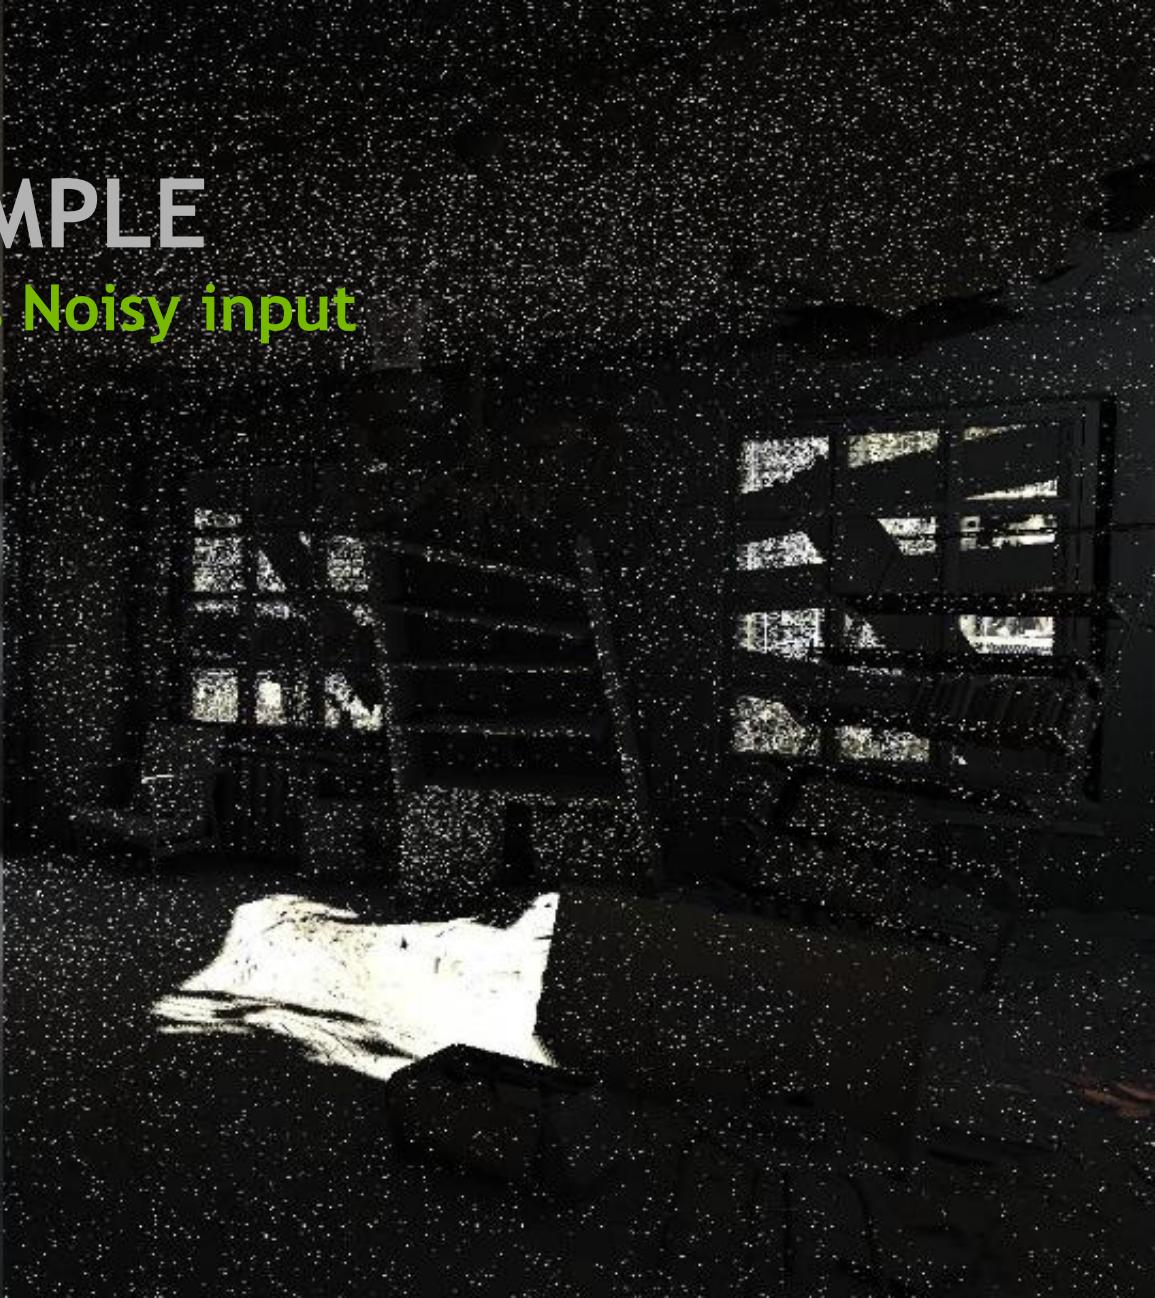
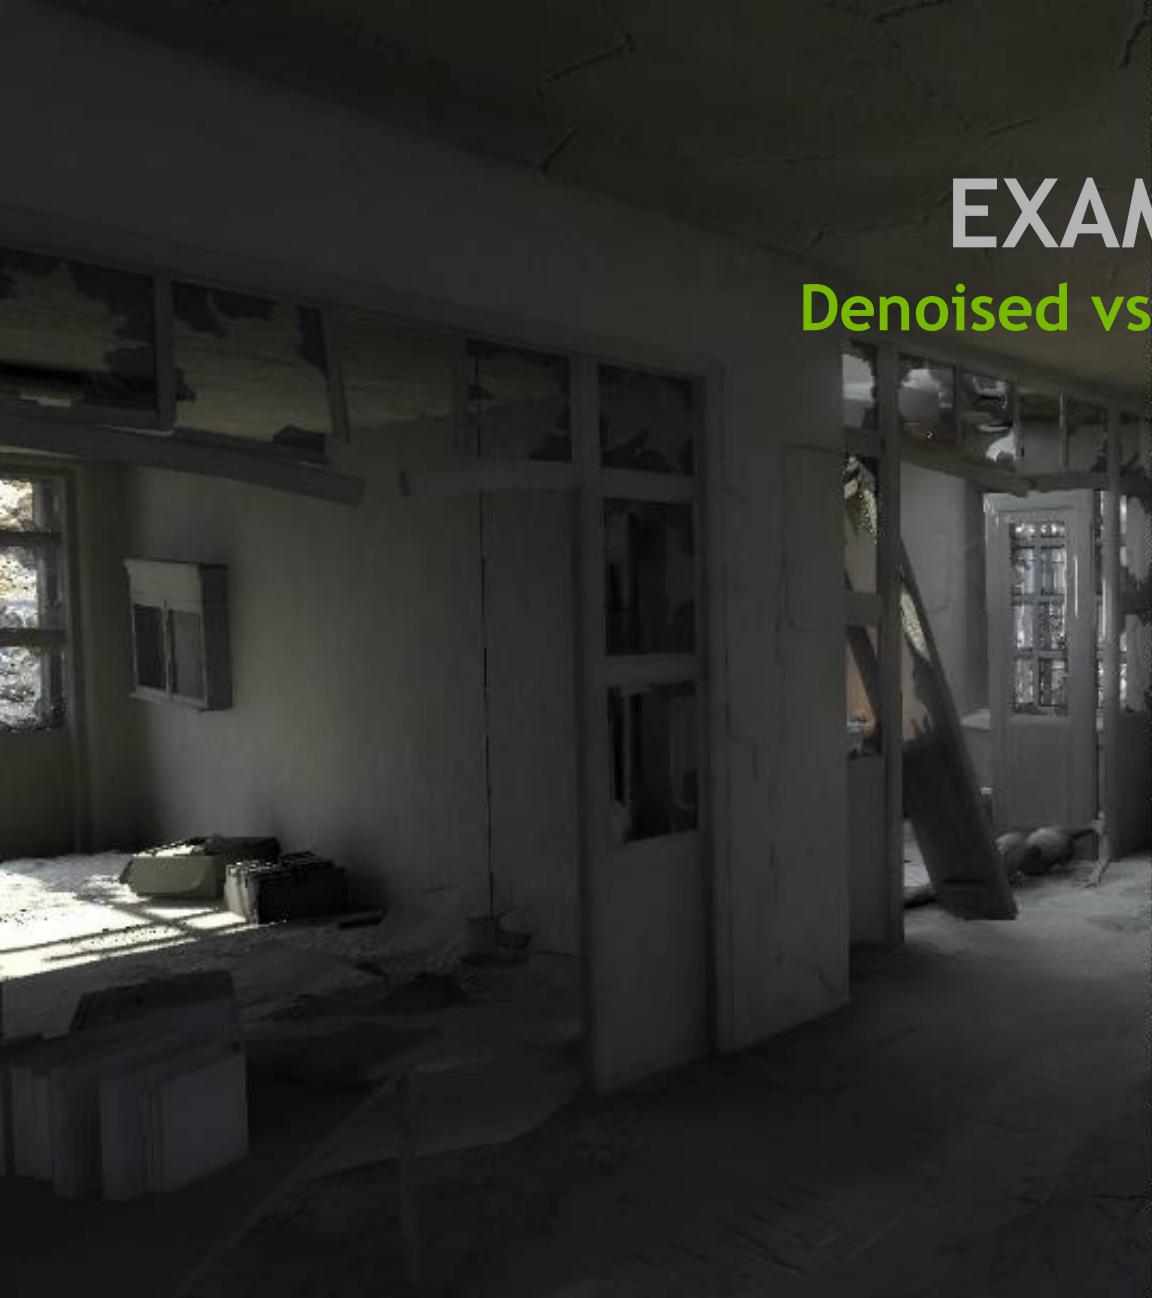
Can be convolution or Deep Learning based

DL-based solution is barely explored in real-time graphics

Our approach is convolution-based and has spatial and temporal components

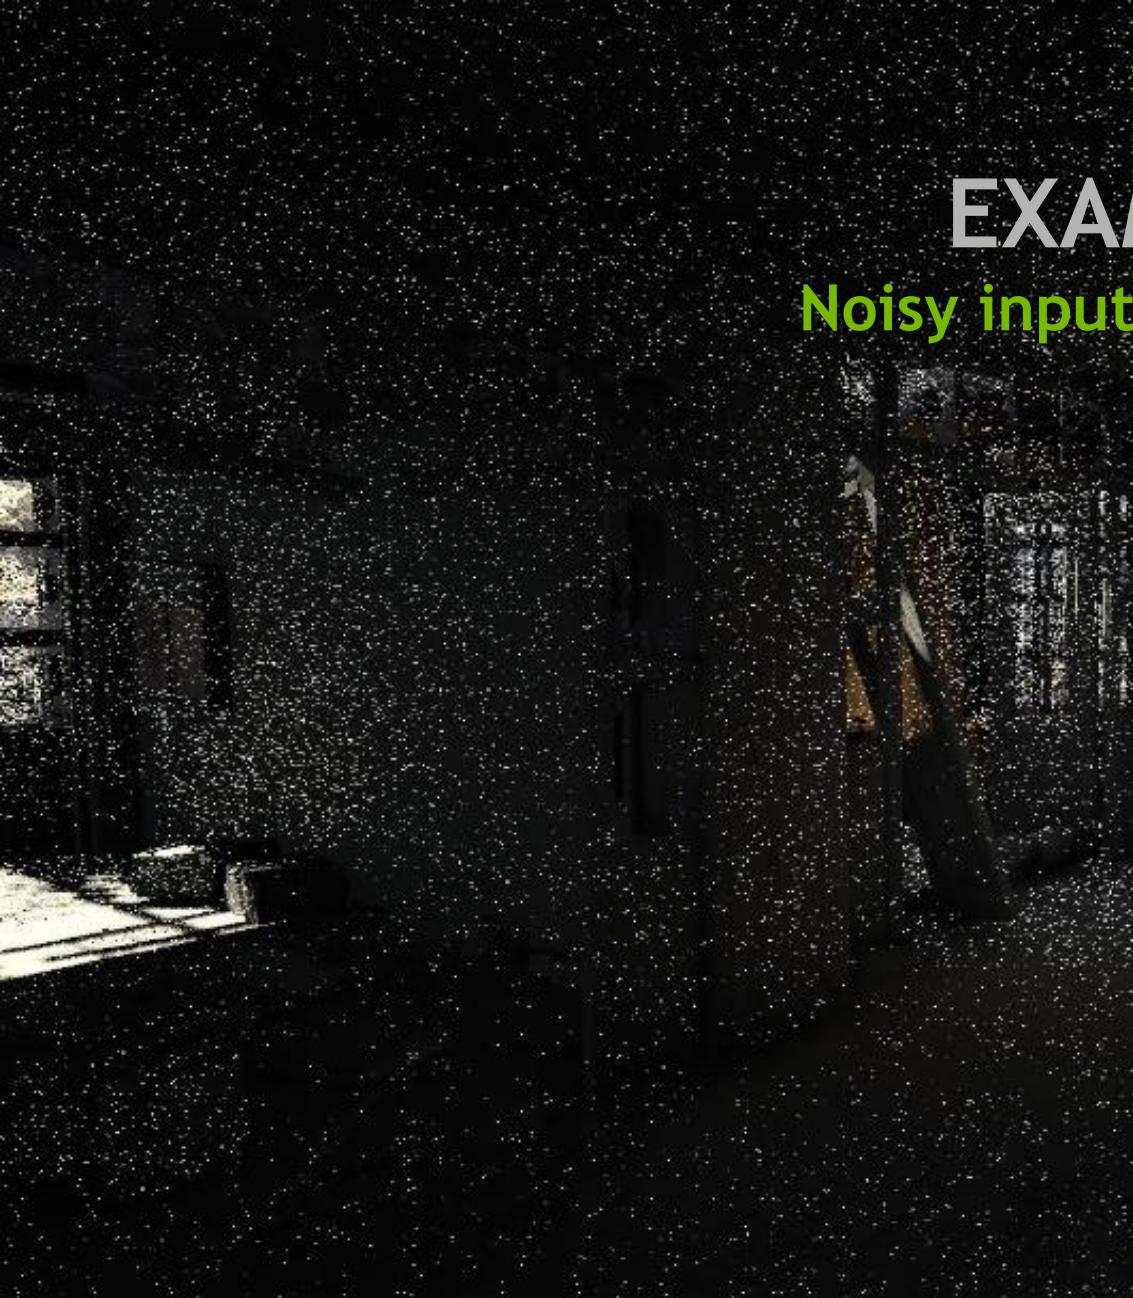
# EXAMPLE

## Denoised vs Noisy input



# EXAMPLE

## Noisy input vs Denoised



# DENOISING IS NOT A FUN...

...but casting rays is :)

Keeps you sad - IQ is always lower than it needs to be

Friendship is very fragile - a small change can ruin IQ completely

Small gifts don't help - tiny tunings here and there turn the algorithm into  
Frankenstein's creature

Demands too much of attention - single pass denoising works badly or inefficiently

# DENOISING

## Problem decomposition

### Spatial component:

Sampling space, distribution and radius?

Sample weight?

Number of samples?

### Temporal component:

Feedback link or links?

Feedback strength and ghosting?

# DENOISING: SPATIAL COMPONENT (1)

As a single-pass blur

Take a lot of samples around current pixel

Accumulate weighted sum

The weight depends on the signal type (AO or GI, reflections, shadows)

Same as Monte Carlo integration:

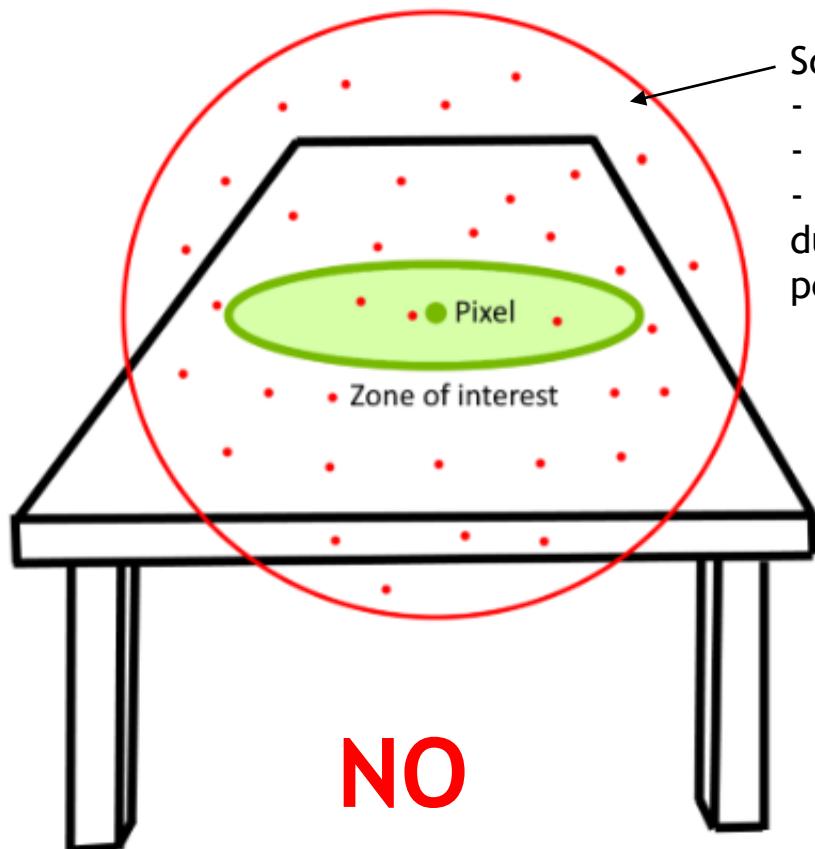
$$\int f(x)dx \approx \frac{1}{N} \sum_{n=1}^N f(x_i)$$

Final reconstructed signal  
(GI, AO)

Weighted sum (N samples)  
 $f(x)$  - noisy input

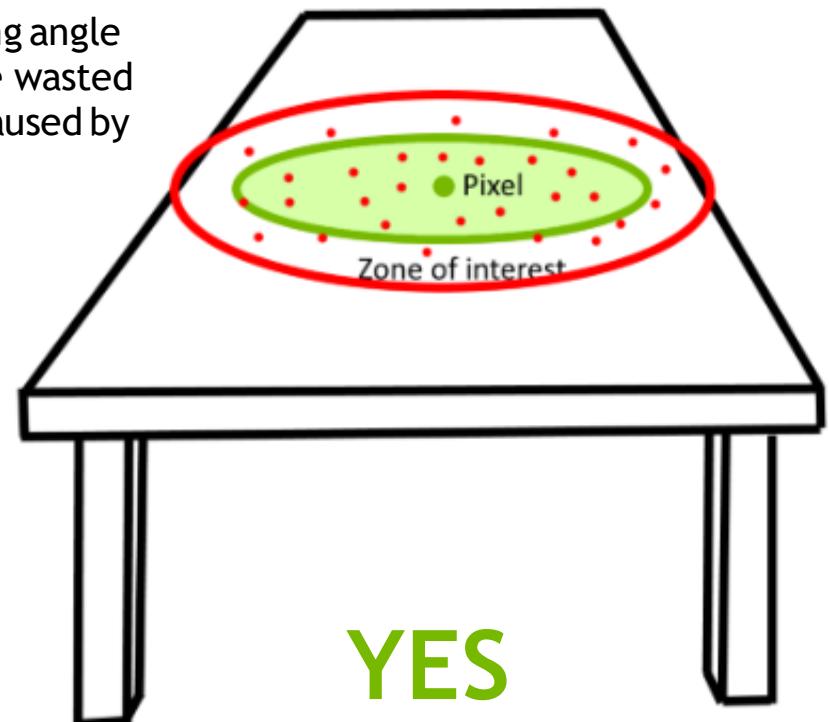
# DENOISING: SPATIAL COMPONENT (2)

## Screen- vs world- space sampling



Screen space problems:

- thin objects
- surfaces at glancing angle
- lots of samples are wasted due to anisotropy caused by perspective



# DENOISING: SPATIAL COMPONENT (3)

## Importance sampling

$$\int f(x)dx \approx \frac{1}{N} \sum_{n=1}^N \frac{f(x_i)}{p(x_i)}$$

Final reconstructed  
signal (GI, AO)

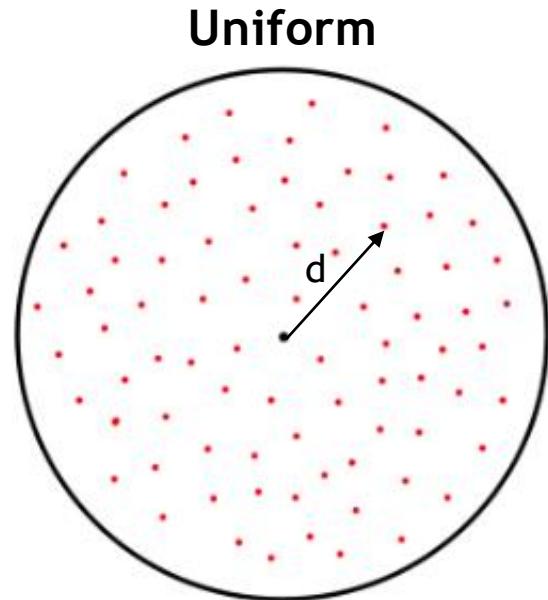
Weighted sum (N samples)  
 $f(x)$  - noisy input

$p(x)$  - Probability Distribution Function (PDF) allows to replace  
uniform distribution with something more relevant...

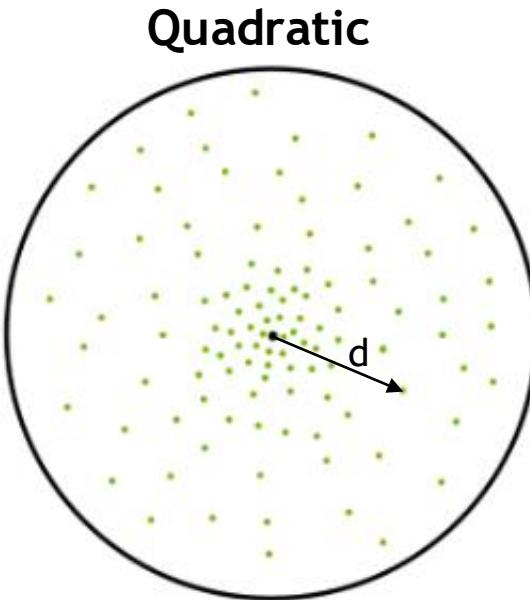
# DENOISING: SPATIAL COMPONENT (4)

Sampling distribution & distance weight

NO



Weight = non\_linear\_F(d)



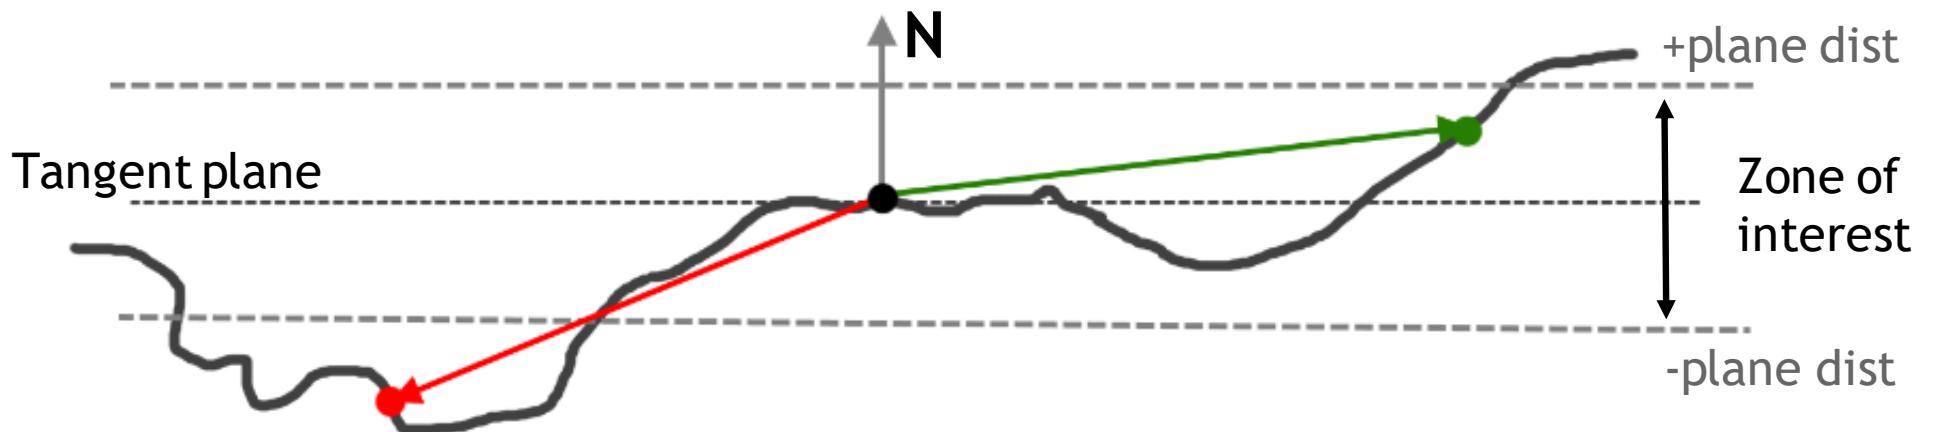
YES

Weight = linear\_F(d) or step(d, R)

Moving distance falloff math to the distribution and simplifying weight calculation to “step” function leads to output noise reduction!

# DENOISING: SPATIAL COMPONENT (5)

Distance weight



Most important samples are on tangent plane

Use plane distance to calculate falloff

Use absolute value, otherwise denoising will skip all rounded objects

# DENOISING: SPATIAL COMPONENT (6)

## Normal weight

Using **pow** is incorrect because it explicitly contradicts lighting theory

It makes your result very oriented

Using **x** instead of **pow(x, 8)** is a good idea

```
// Please, don't use 'pow' !
float NormalWeight(float3 Ncenter, float3 Nsample)
{
    float f = dot(Ncenter, Nsample);

    return pow(saturate(f), 8.0);
}
```

# DENOISING: SPATIAL COMPONENT (7)

Per pixel kernel rotations

**NO!**

Leads to 2x-5x slowdown!

Input signal is already noisy (applying noise on top of noise isn't worth it)

Use **per frame random rotation** to improve quality of temporal accumulation!

# DENOISING: SPATIAL COMPONENT (8)

## Radius of denoising

Needs to be large, but can be scaled with distance

Compute **variance** of the input signal, blur less if variance is small

Blur less in “dark corners”, i.e. multiply by AO

**Signal-to-noise ratio** - blur less where direct lighting is strong

$$R = \text{BaseRadius} \cdot F(\text{viewZ}) \cdot F(\text{variance}) \cdot F(\text{AO})$$

# DENOISING: SPATIAL COMPONENT (8)

## Number of samples

A lot of samples are required! 32? 64? 128? (depending on number of passes)

Compute **variance** of the input signal, adaptively reduce number of samples if variance of the input signal is small...

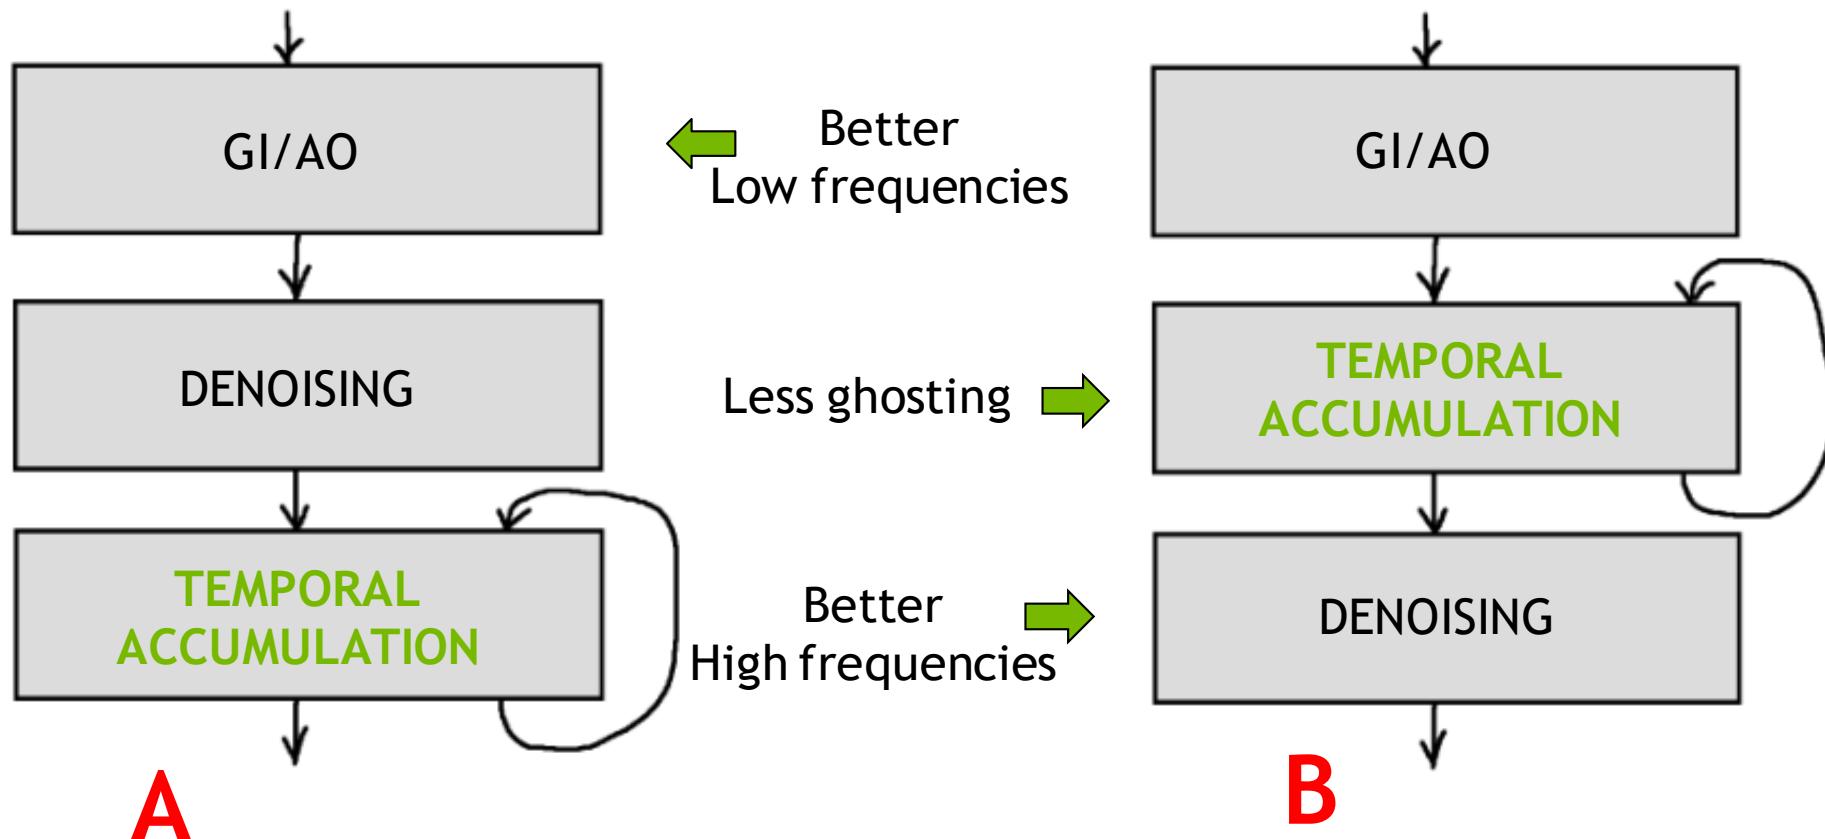
...but variance computed for the current frame is always big! **Solution - add temporal component \O/**

Obviously, accumulated signal will get less and less variance over time!



# DENOISING: TEMPORAL COMPONENT (1)

Common ideas



# DENOISING: TEMPORAL COMPONENT (2)

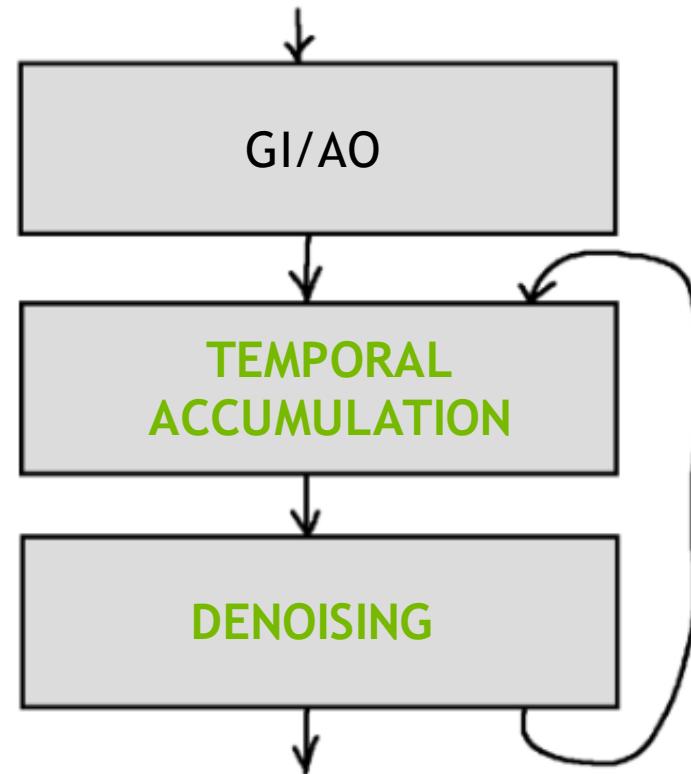
Our idea

More frequencies over time (mixture of low and high)

Requires less samples per frame

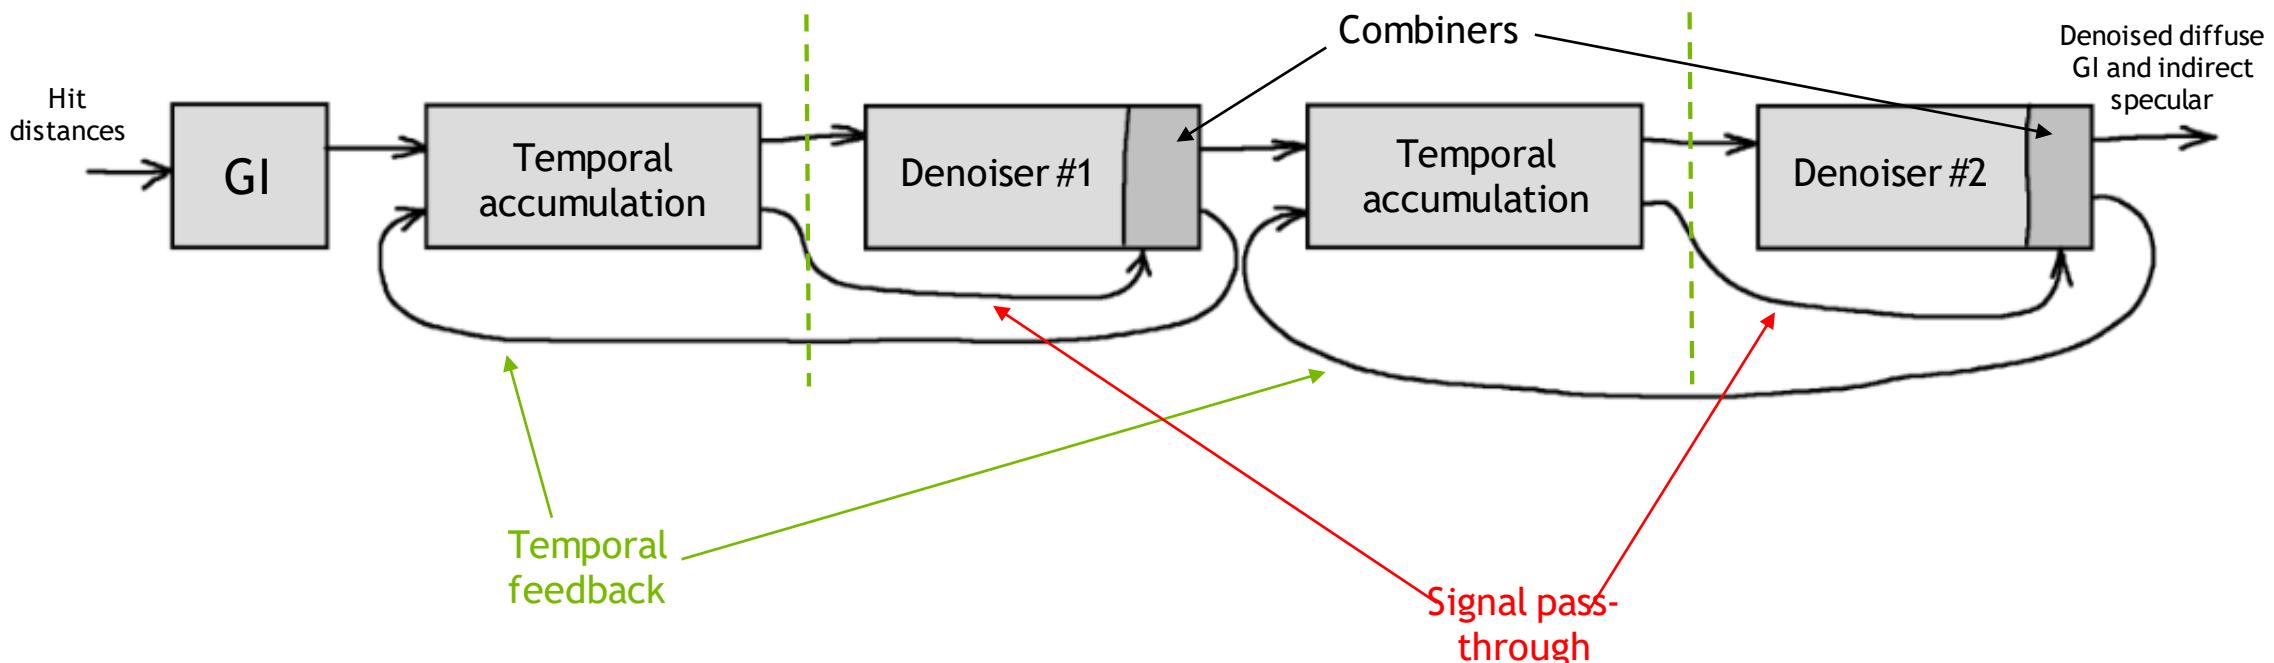
Less ghosting (denoising smoothes out reprojection artefacts)

(AO denoising uses this scheme, adaptive sampling with up to 64 samples, processes 2 pixels per thread sharing results between them if no edges)



# DENOISING: LITTLE MONSTER (1)

## GI denoiser



# DENOISING: LITTLE MONSTER (2)

## Denoiser block

Computes variance of the input signal (3x3 pixels)

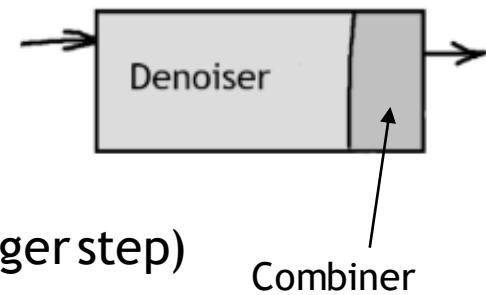
Computes radius scale as “ $F(\text{viewZ}) \cdot F(\text{variance}) \cdot F(\text{AO})$ ”

Computes adaptive step  $N = F(\text{scaleRadius})$  (small radius = bigger step)

Processes each Nth sample from a poisson disk (up to 32 samples per pass)

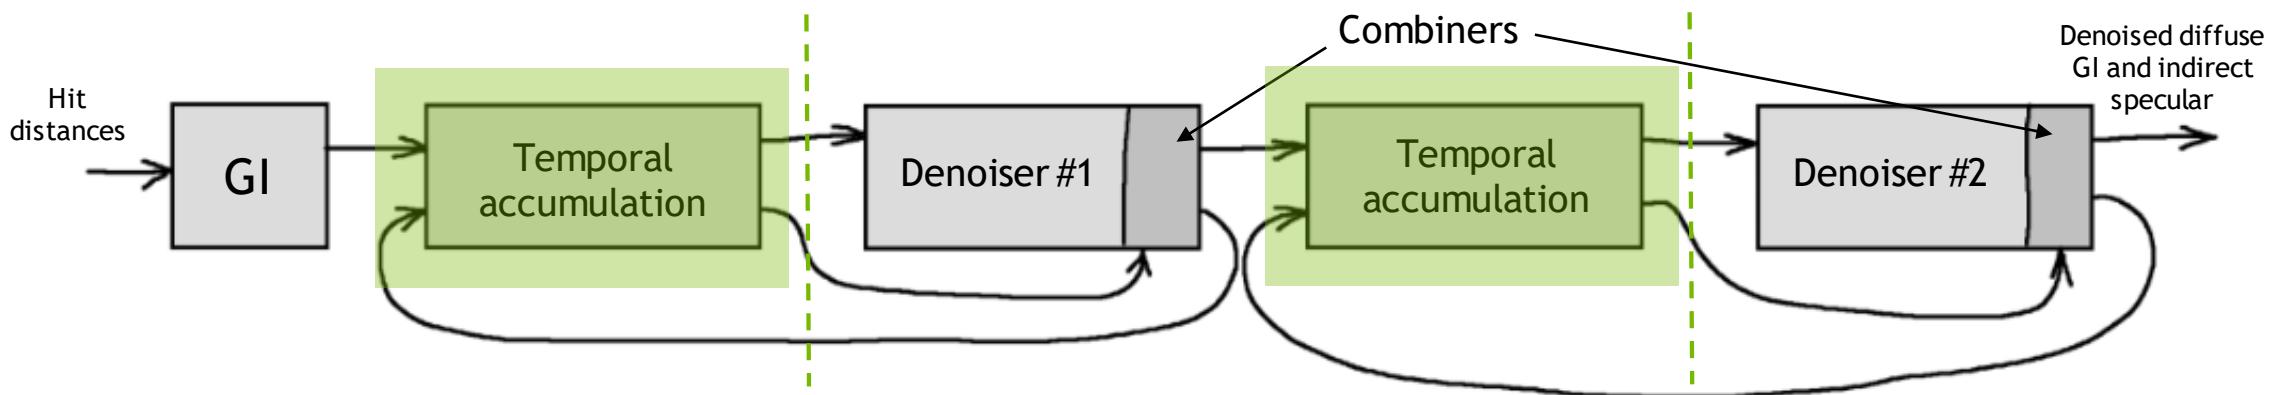
The combiner just mixes up denoised and noisy input signals as:

```
Combiner = lerp(denoisedSignal, inputSignal, 0.5 * accumSpeed) (accumSpeed = 0.93 if no motion)
```



# DENOISING: LITTLE MONSTER (3)

## GI denoiser

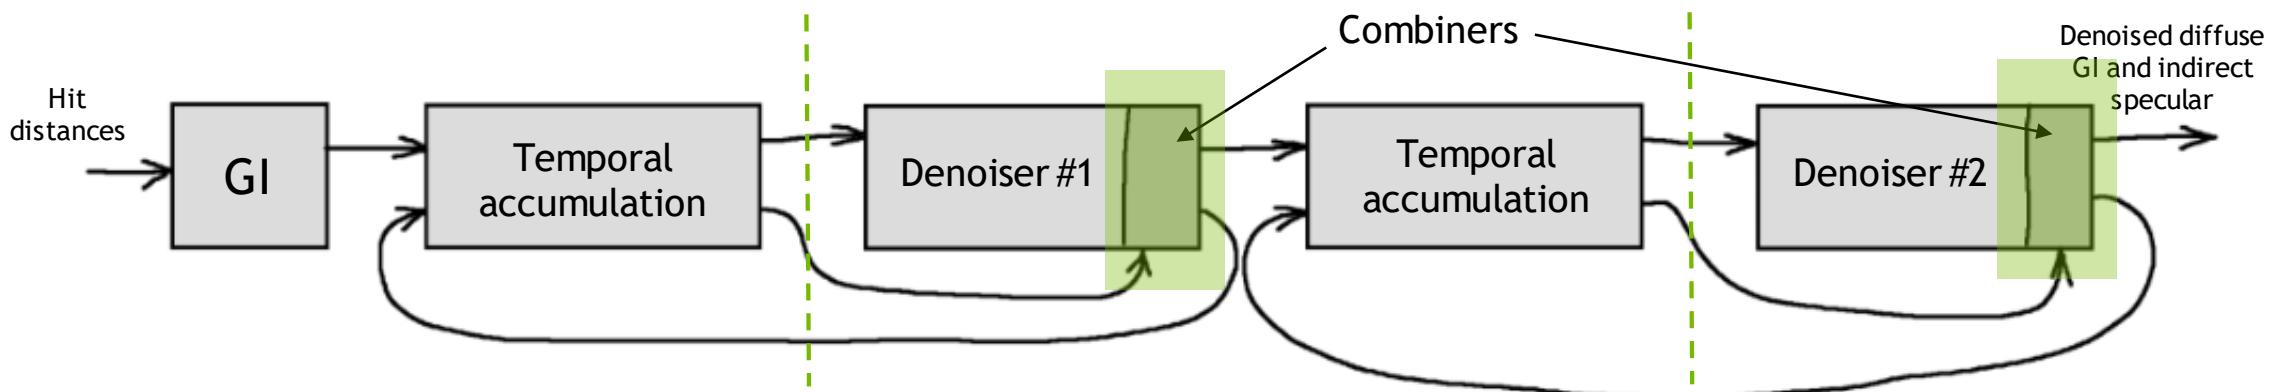


**Temporal accumulation** always happens **before denoising** to eliminate ghosting and reprojection artefacts

History is always rejected if out-of-screen sampling or z-occlusion are detected

# DENOISING: LITTLE MONSTER (4)

## GI denoiser

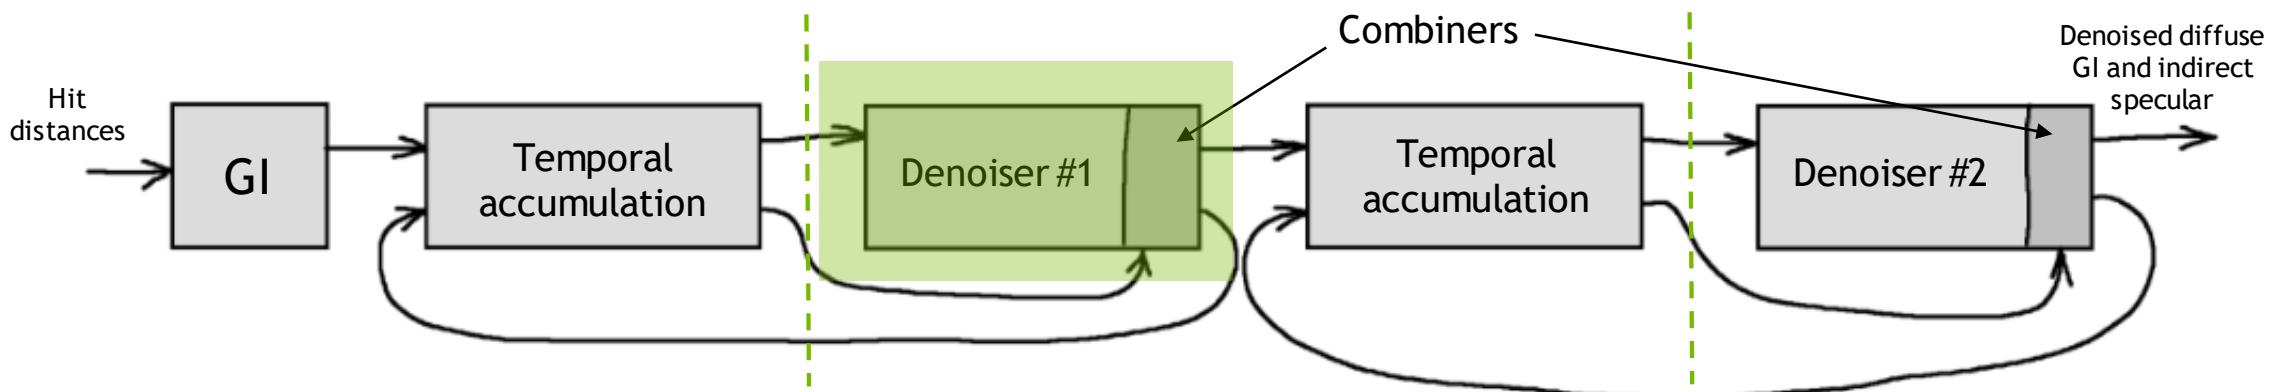


The **output** of each denoiser is always a **combination of denoised and noisy input signals!**

It helps to preserve tiny details

# DENOISING: LITTLE MONSTER (5)

GI denoiser

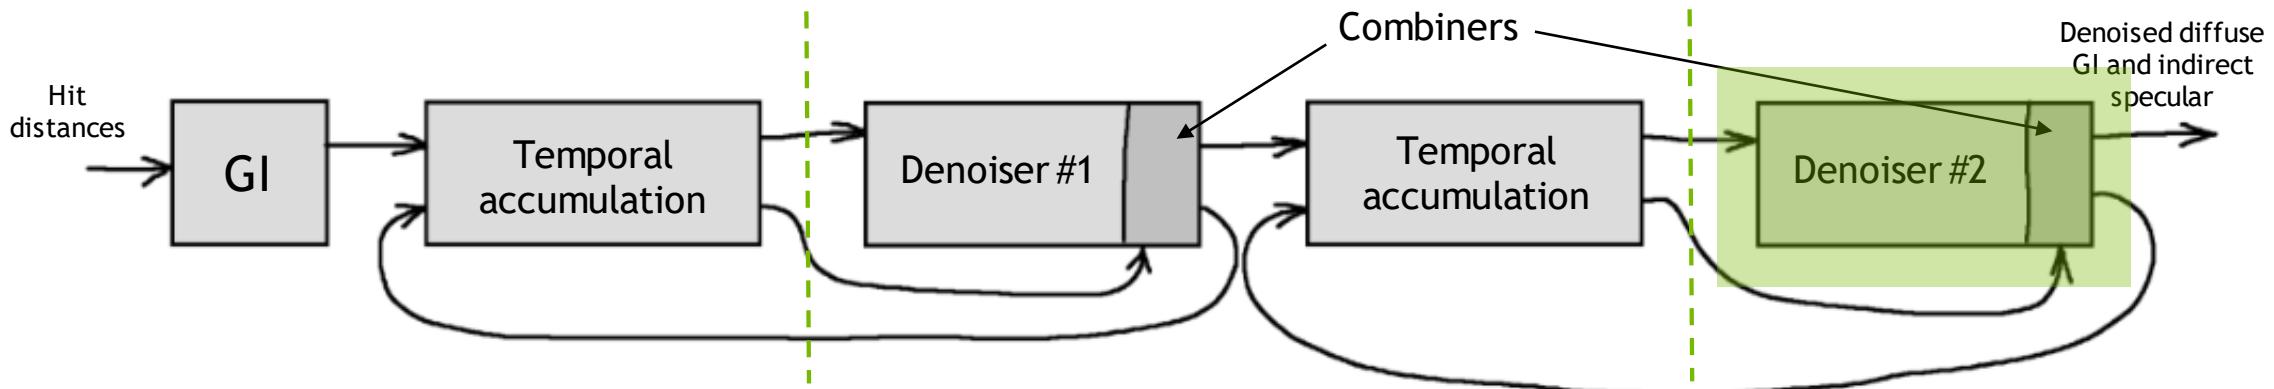


First pass of denoising doesn't take normals into account

It has wider base radius (6m)

# DENOISING: LITTLE MONSTER (6)

GI denoiser



Second pass of denoising takes normals into account

It has smaller base radius (3m)

Physically it's same denoiser which applies “normal weight” on top of geometry weight

# DENOISING

## Tips & tricks

Use **NSIGHT GRAPHICS** GPU Trace utility to understand your limiters

**Fetch** heavy data only **if weight is non-zero**

**TAA** is your friend - it's a free pass of denoising

**SH irradiance** is your friend - solves “blurriness” problem

**Know your noise** - perfection in image “cleanliness” is not needed

# PERFORMANCE

## RTX 2080 at 2560x1440

Stage	HIGH	ULTRA
Pretrace	~0.4 ms	~0.8 ms
BLAS/TLAS (completely hidden by async)	~0.5 ms	~0.5 ms
Raytracing	1 to 3 ms	2 to 6 ms
AO Denoising	~0.6 ms	~0.9 ms
GI computation	~0.6 ms	~1.0 ms
GI denoising	~1.6 ms	~2.1 ms
<b>Total Frame Time Overhead (vs RTX OFF)</b>	<b>~20%</b>	<b>~30%</b>



# ARTIST POINT OF VIEW

Just make it work for us



# OUR FIRST RTAO SHOT

...Which one is RT ON? 🤔



# DEFENDING THE CHOICE OF RTGI

Why not do reflections instead?

There were not many people who believed RTGI was a good direction of research

From audience to stakeholders (oops)

Especially when convincing solutions already exist:

- SSAO and geometric ESM-AO for world space AO

- Super-lazy-realtime grid of probes for GI

- Voxel GI (which we already have nicely integrated with PBR in Exodus)

# LIMITATIONS ARE ALSO WELL KNOWN

And we accepted them for years

Reflection probes or lightmaps for GI?

not a realtime solution

SSAO for AO?

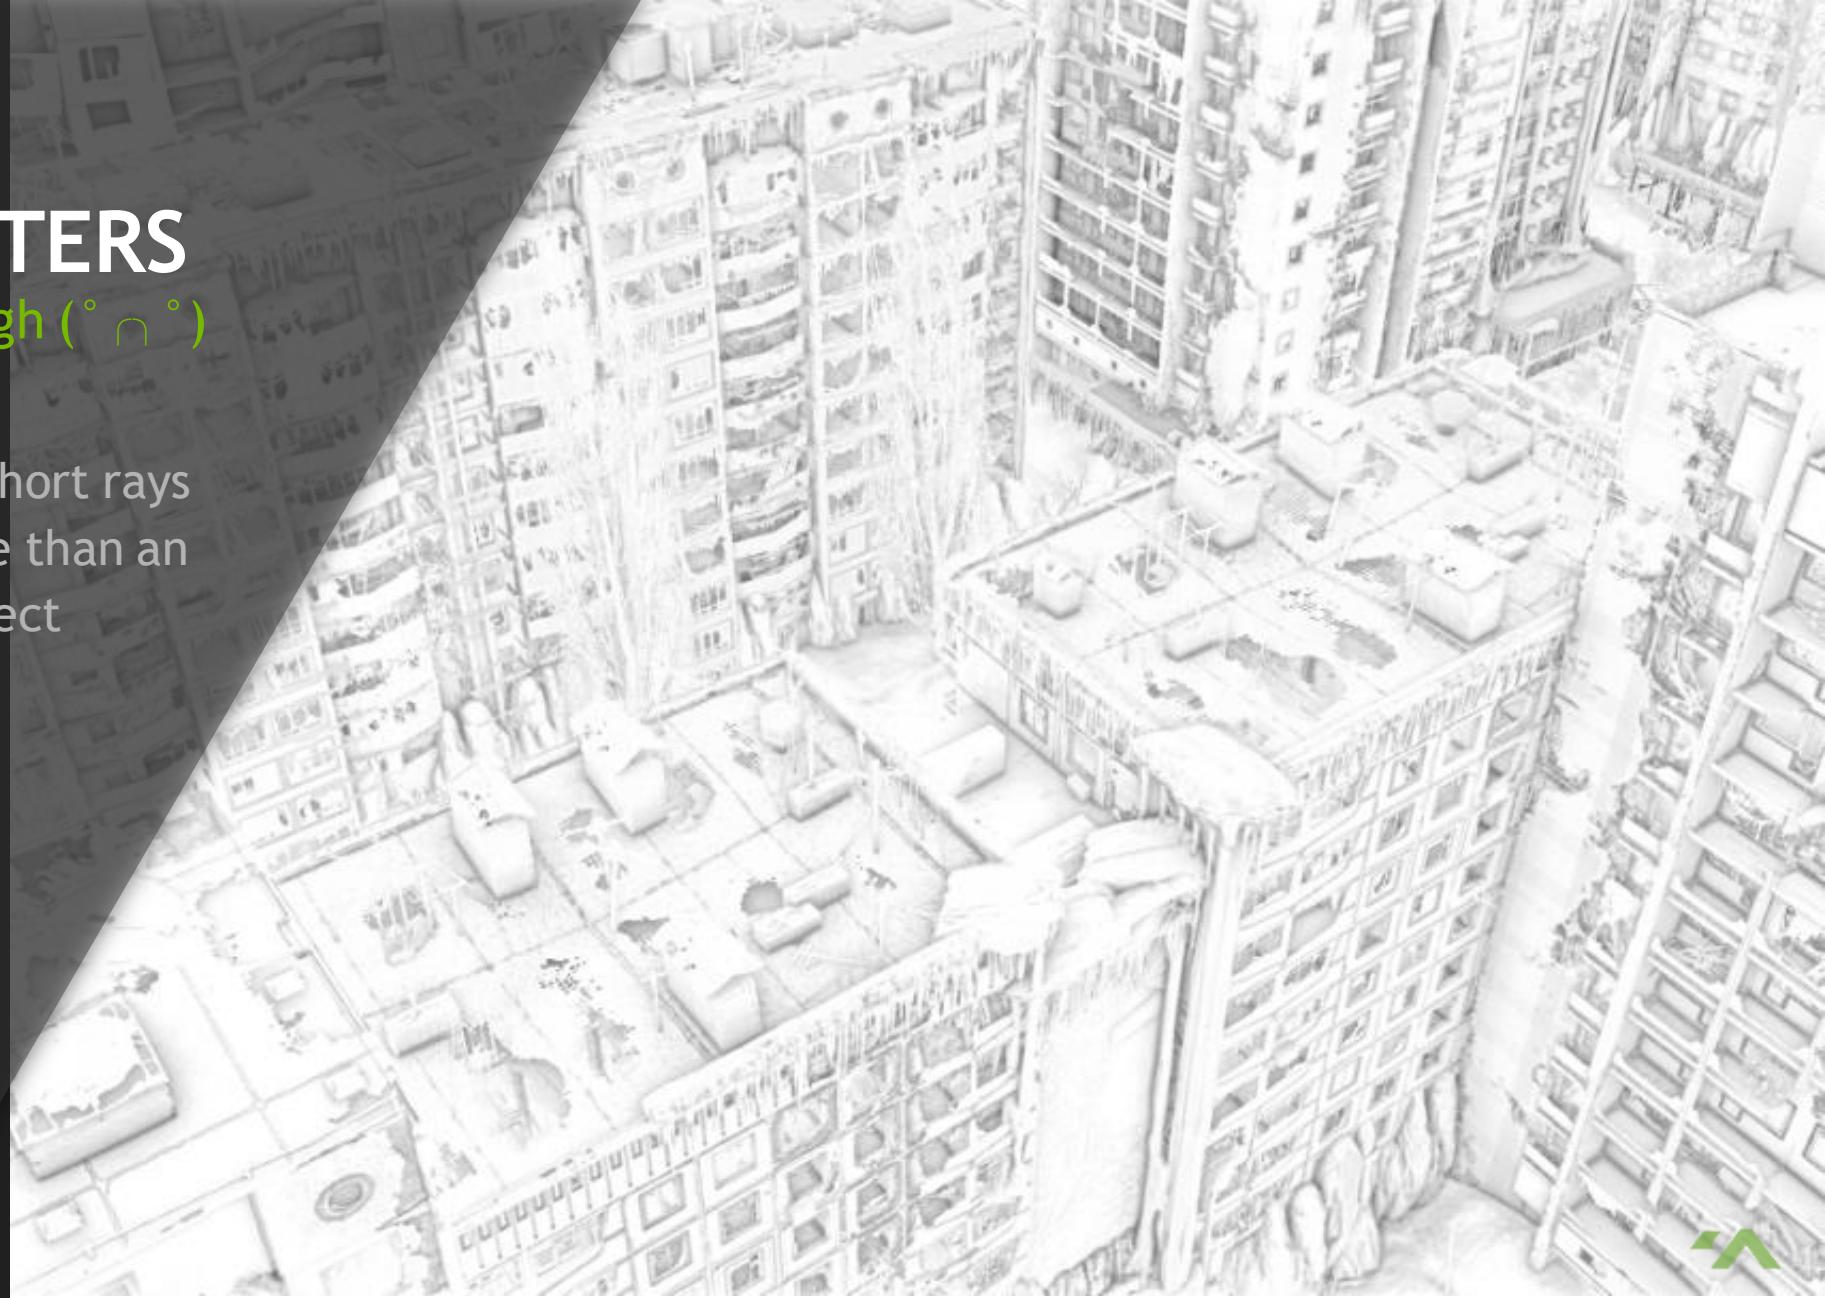
suffers from its screen-space nature

limited to 1m tracing (good for features of... <1m in size)

# SIZE MATTERS

1m is not enough (° ՞ °)

In large scenes short rays  
produce no more than an  
'edge trace' effect



# NEW INSANE POSSIBILITIES

Literally insane

50m ray tracing

Billions of rays per second

Per-pixel details at any scale:

pencils on table

1mm scale

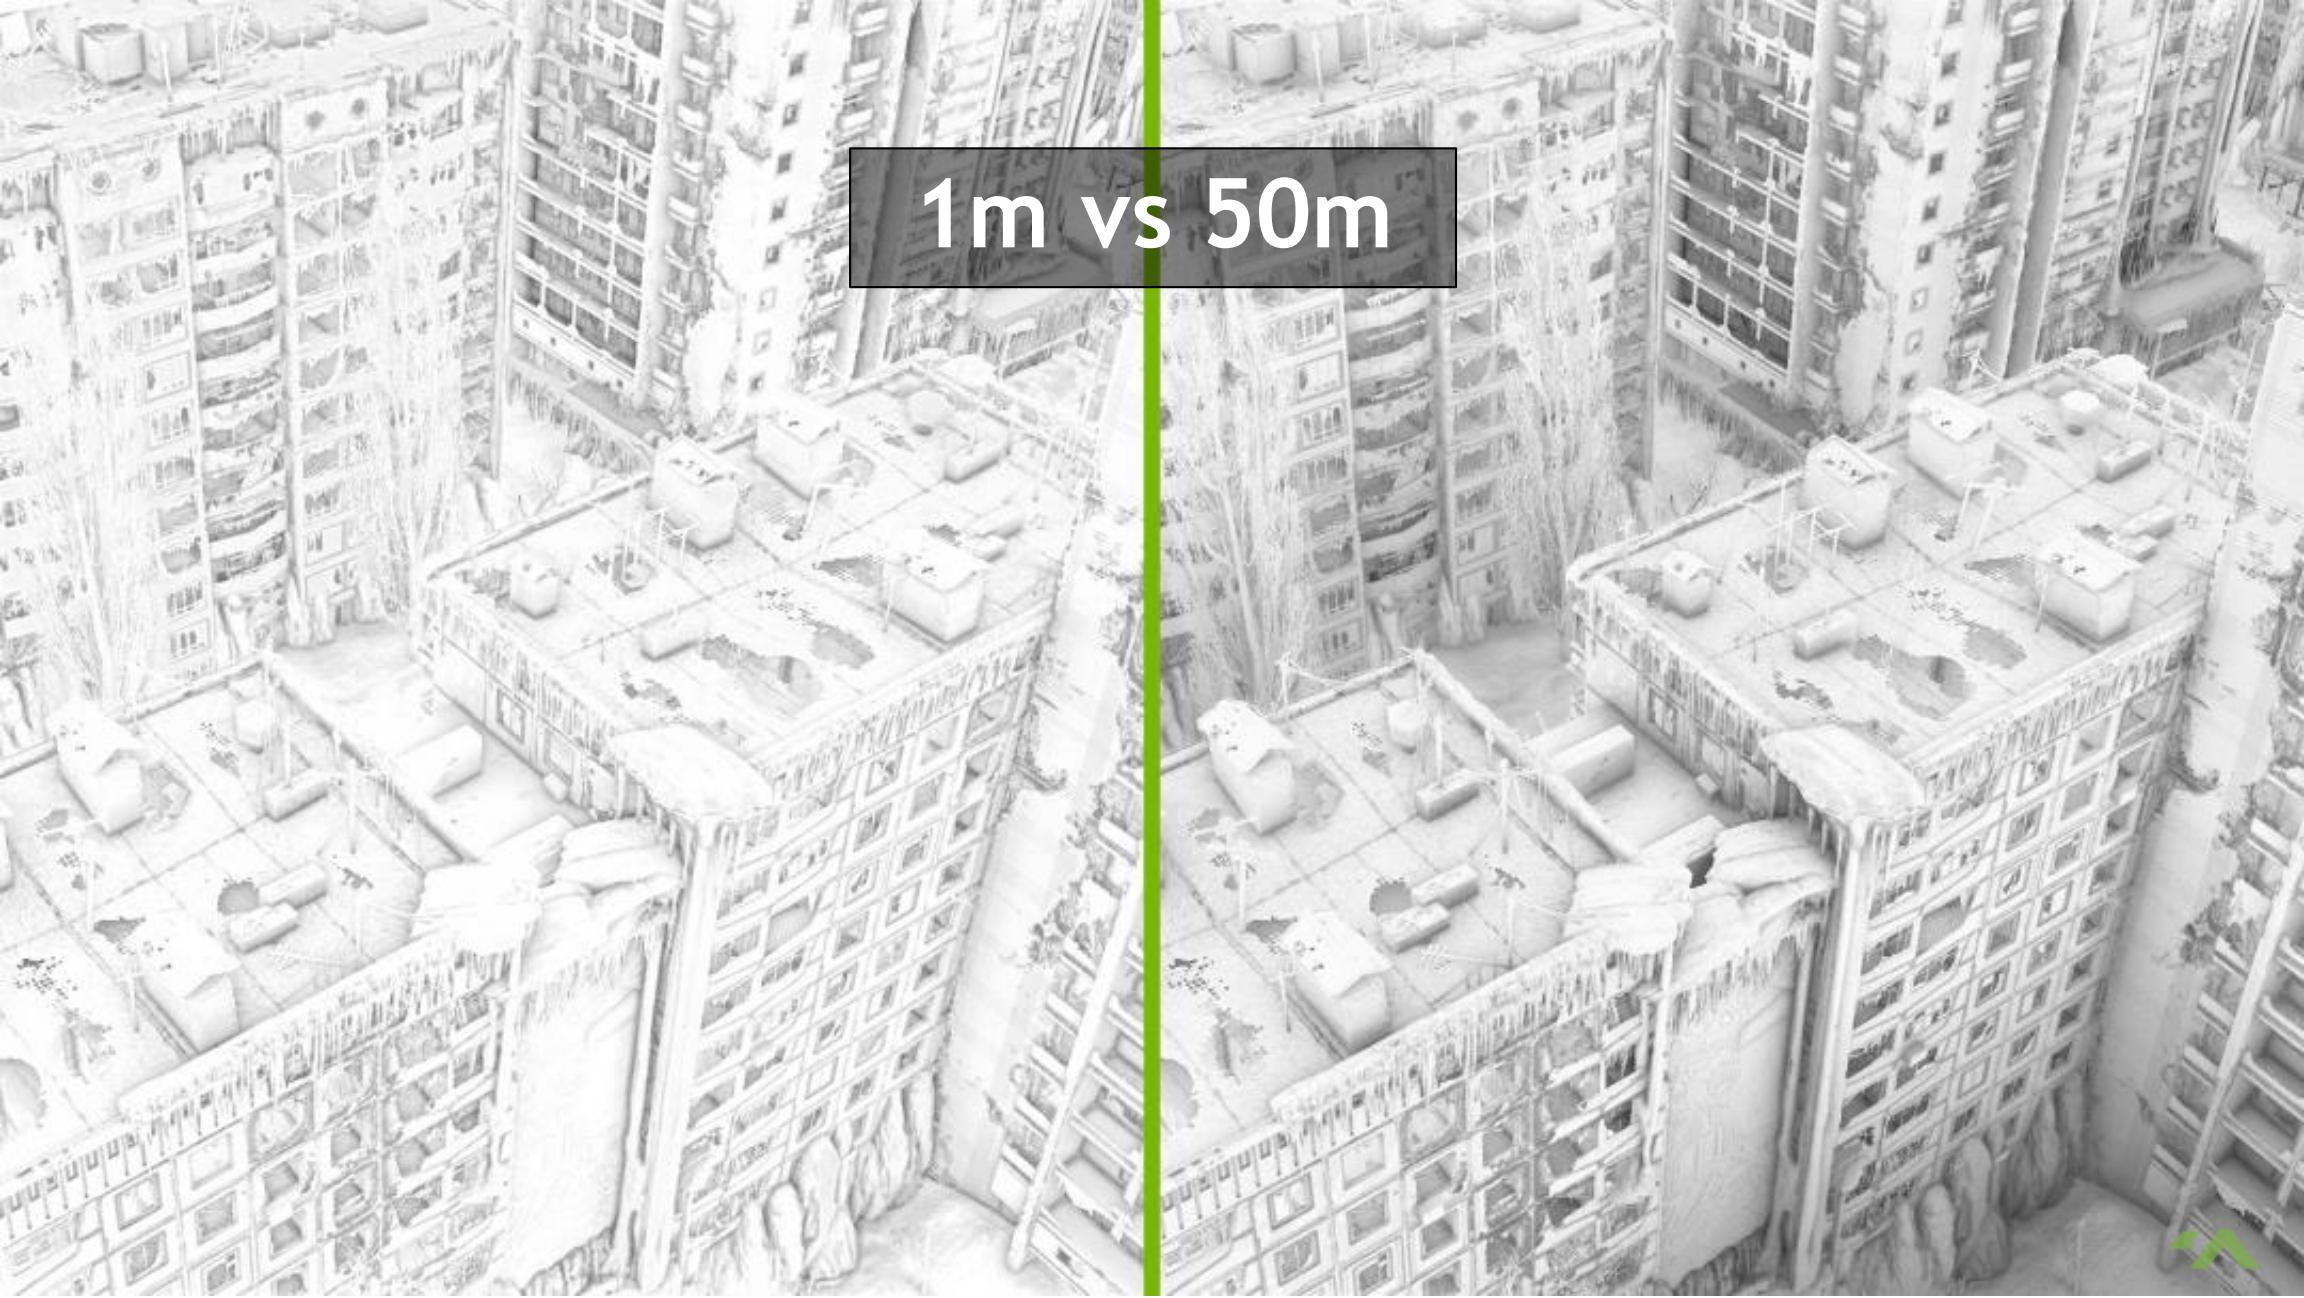
ships

20m scale

canyons, skyscrapers

100m+ scale

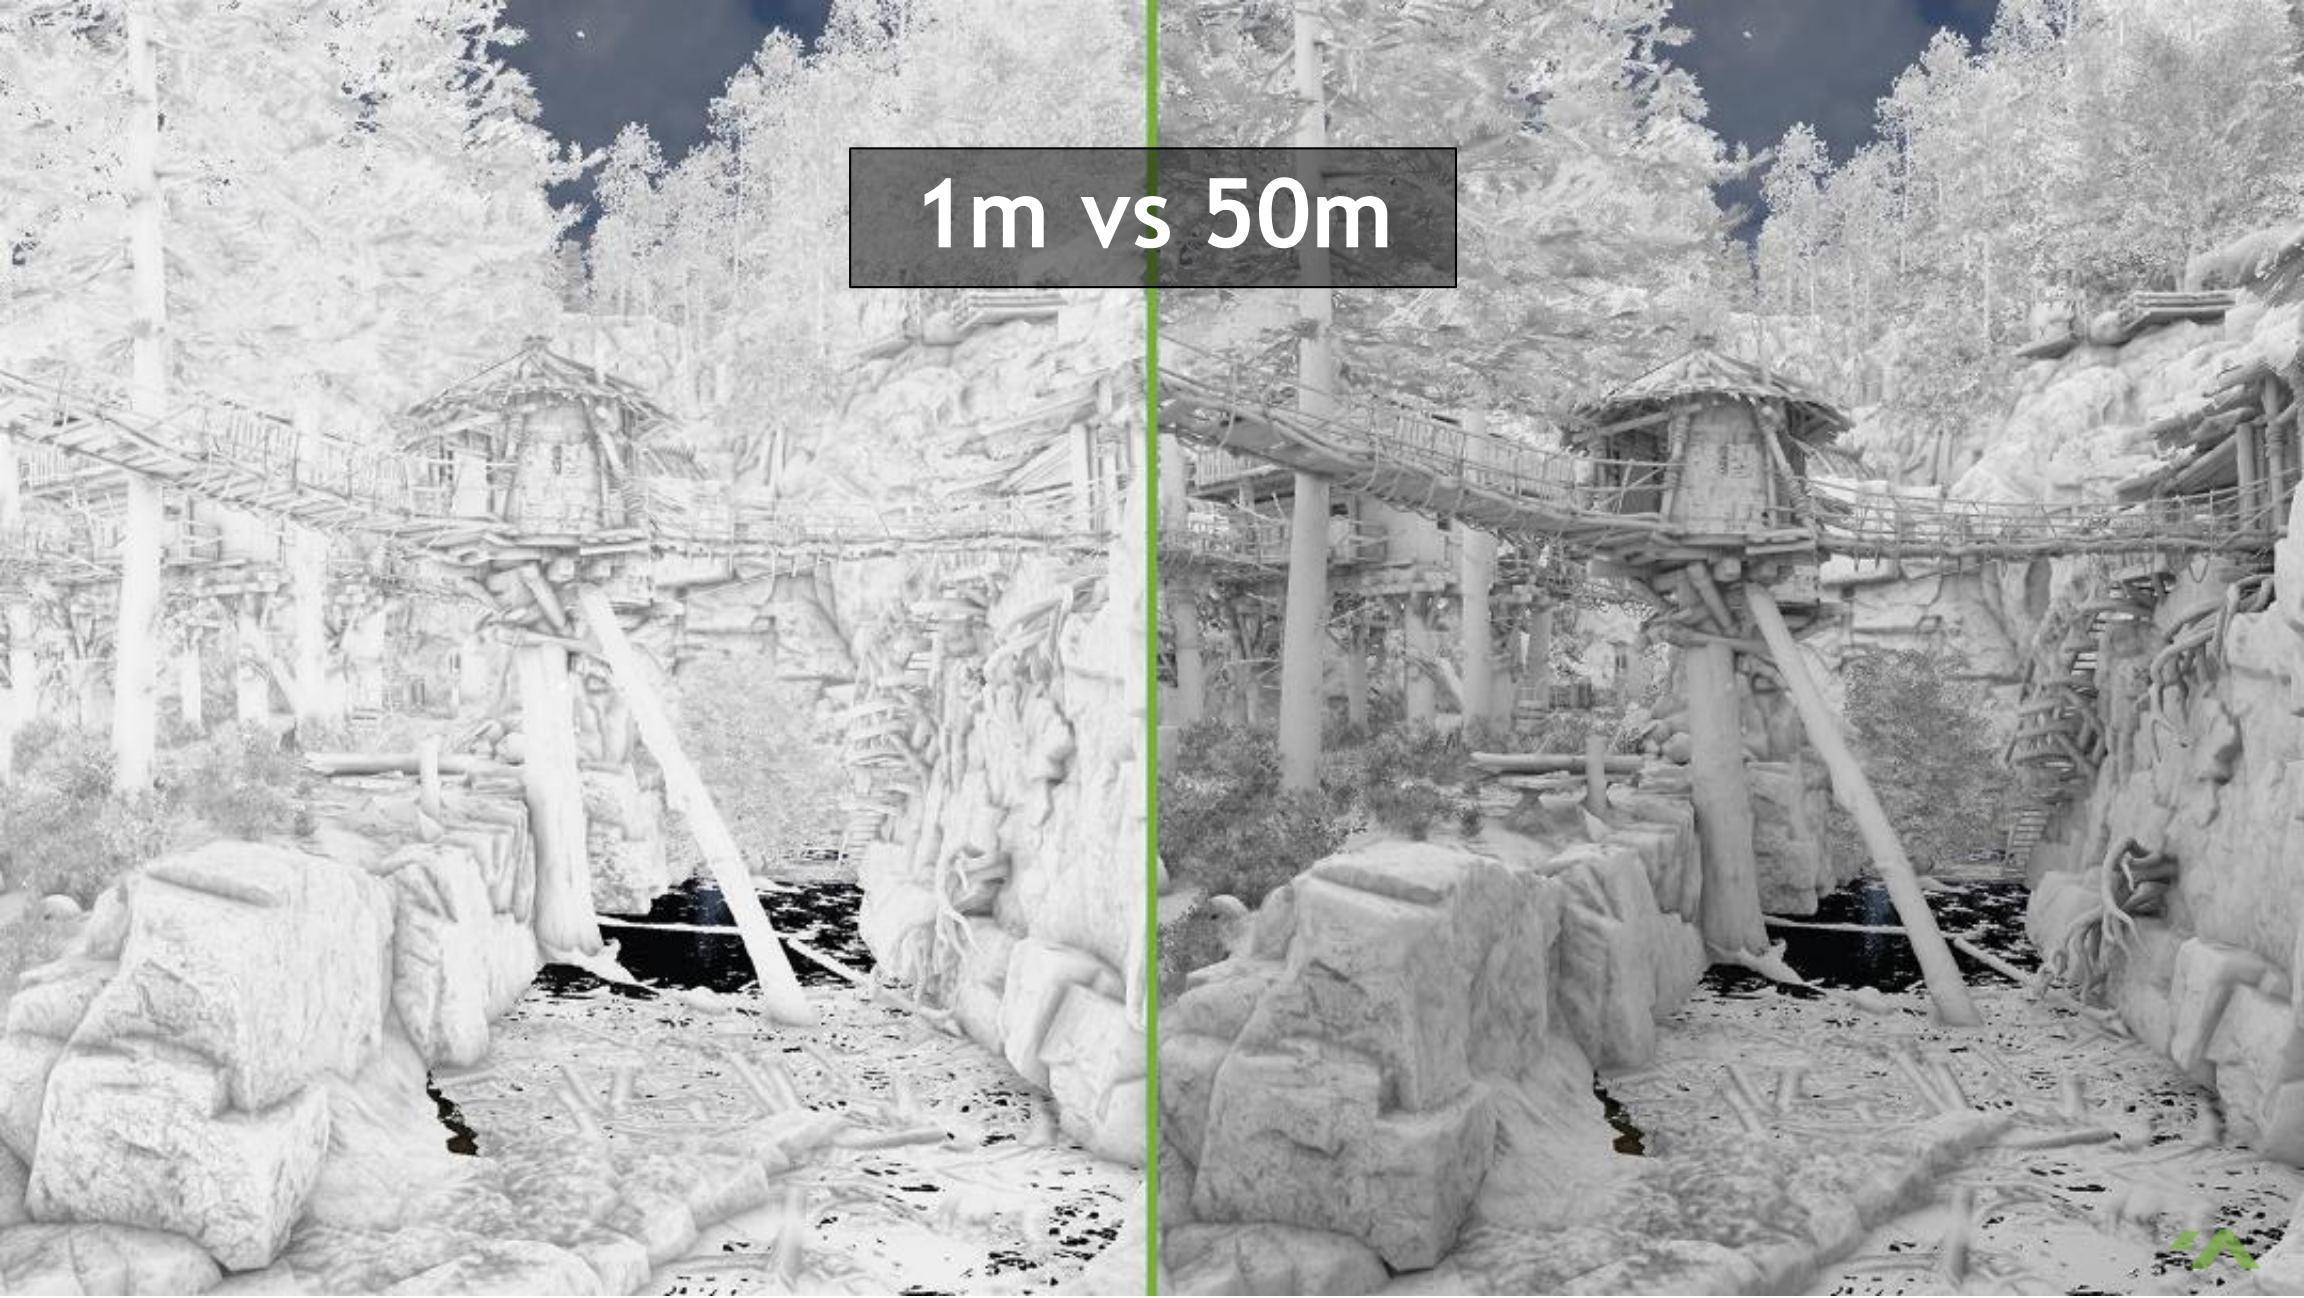
And at no cost!.. Well, almost



A grayscale aerial photograph of a dense urban residential area, likely Hong Kong, featuring numerous high-rise apartment buildings packed closely together. The buildings have numerous windows and some external laundry hanging. In the center of the image, there is a vertical green line that divides the scene into two halves. A black rectangular box is positioned in the upper left quadrant of the image, containing the text "1m vs 50m".

1m vs 50m





**1m vs 50m**





LEGACY AO



RTGI



# SSAO NO MORE

GI replaces the need for it

Legacy AO:

Tons of AO sources mixed

Multiplied directly on shadows

Effectively a patch

RTGI:

Solves it all



# SKYLIGHT SHADOWS

No direct lights involved

Single frame took several  
minutes of rendering in '99

Mesmerizing to watch



# GI FROM LIGHT SOURCES

Interiors fully lit by sun

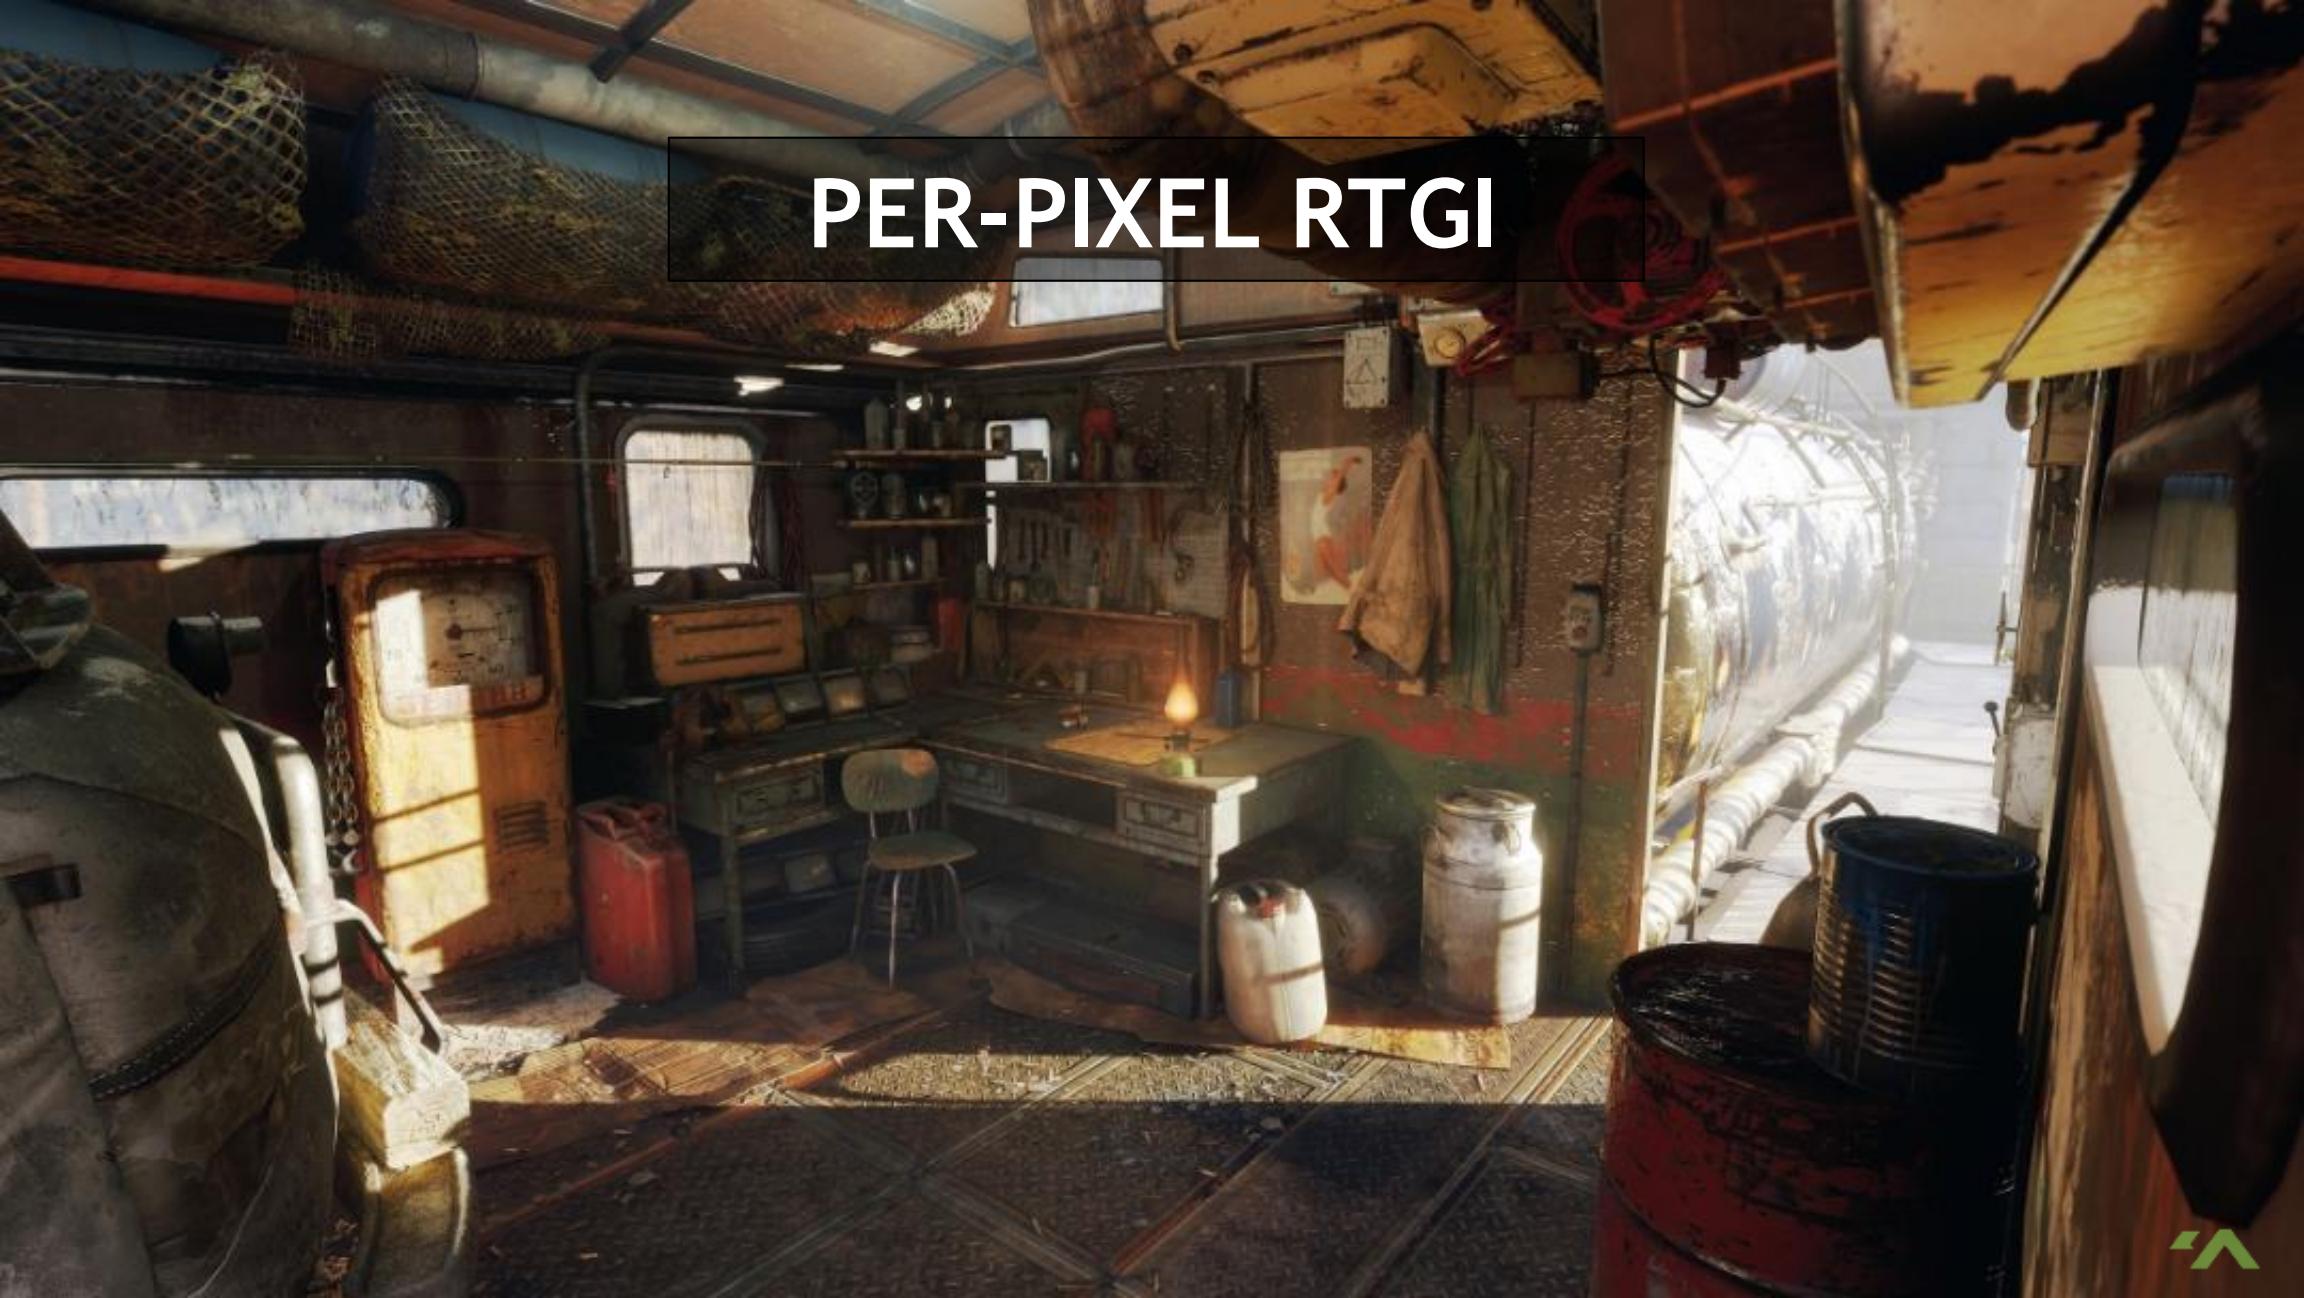
Пиши умное, э



# GI BY LIGHT PROBES



# PER-PIXEL RTGI



# IMPLEMENTATION CONTINUES...

Still missing something

## Specular GI

Specular lighting contributes up to 50% of light on rough surfaces

## Color bleeding

The most prominent feature in GI



# COLOR BLEEDING OFF



# COLOR BLEEDING ON



# CLOSE TO RELEASE

## Content fixes and polishing

Making content work well in both modes

Revert fake artsy lights

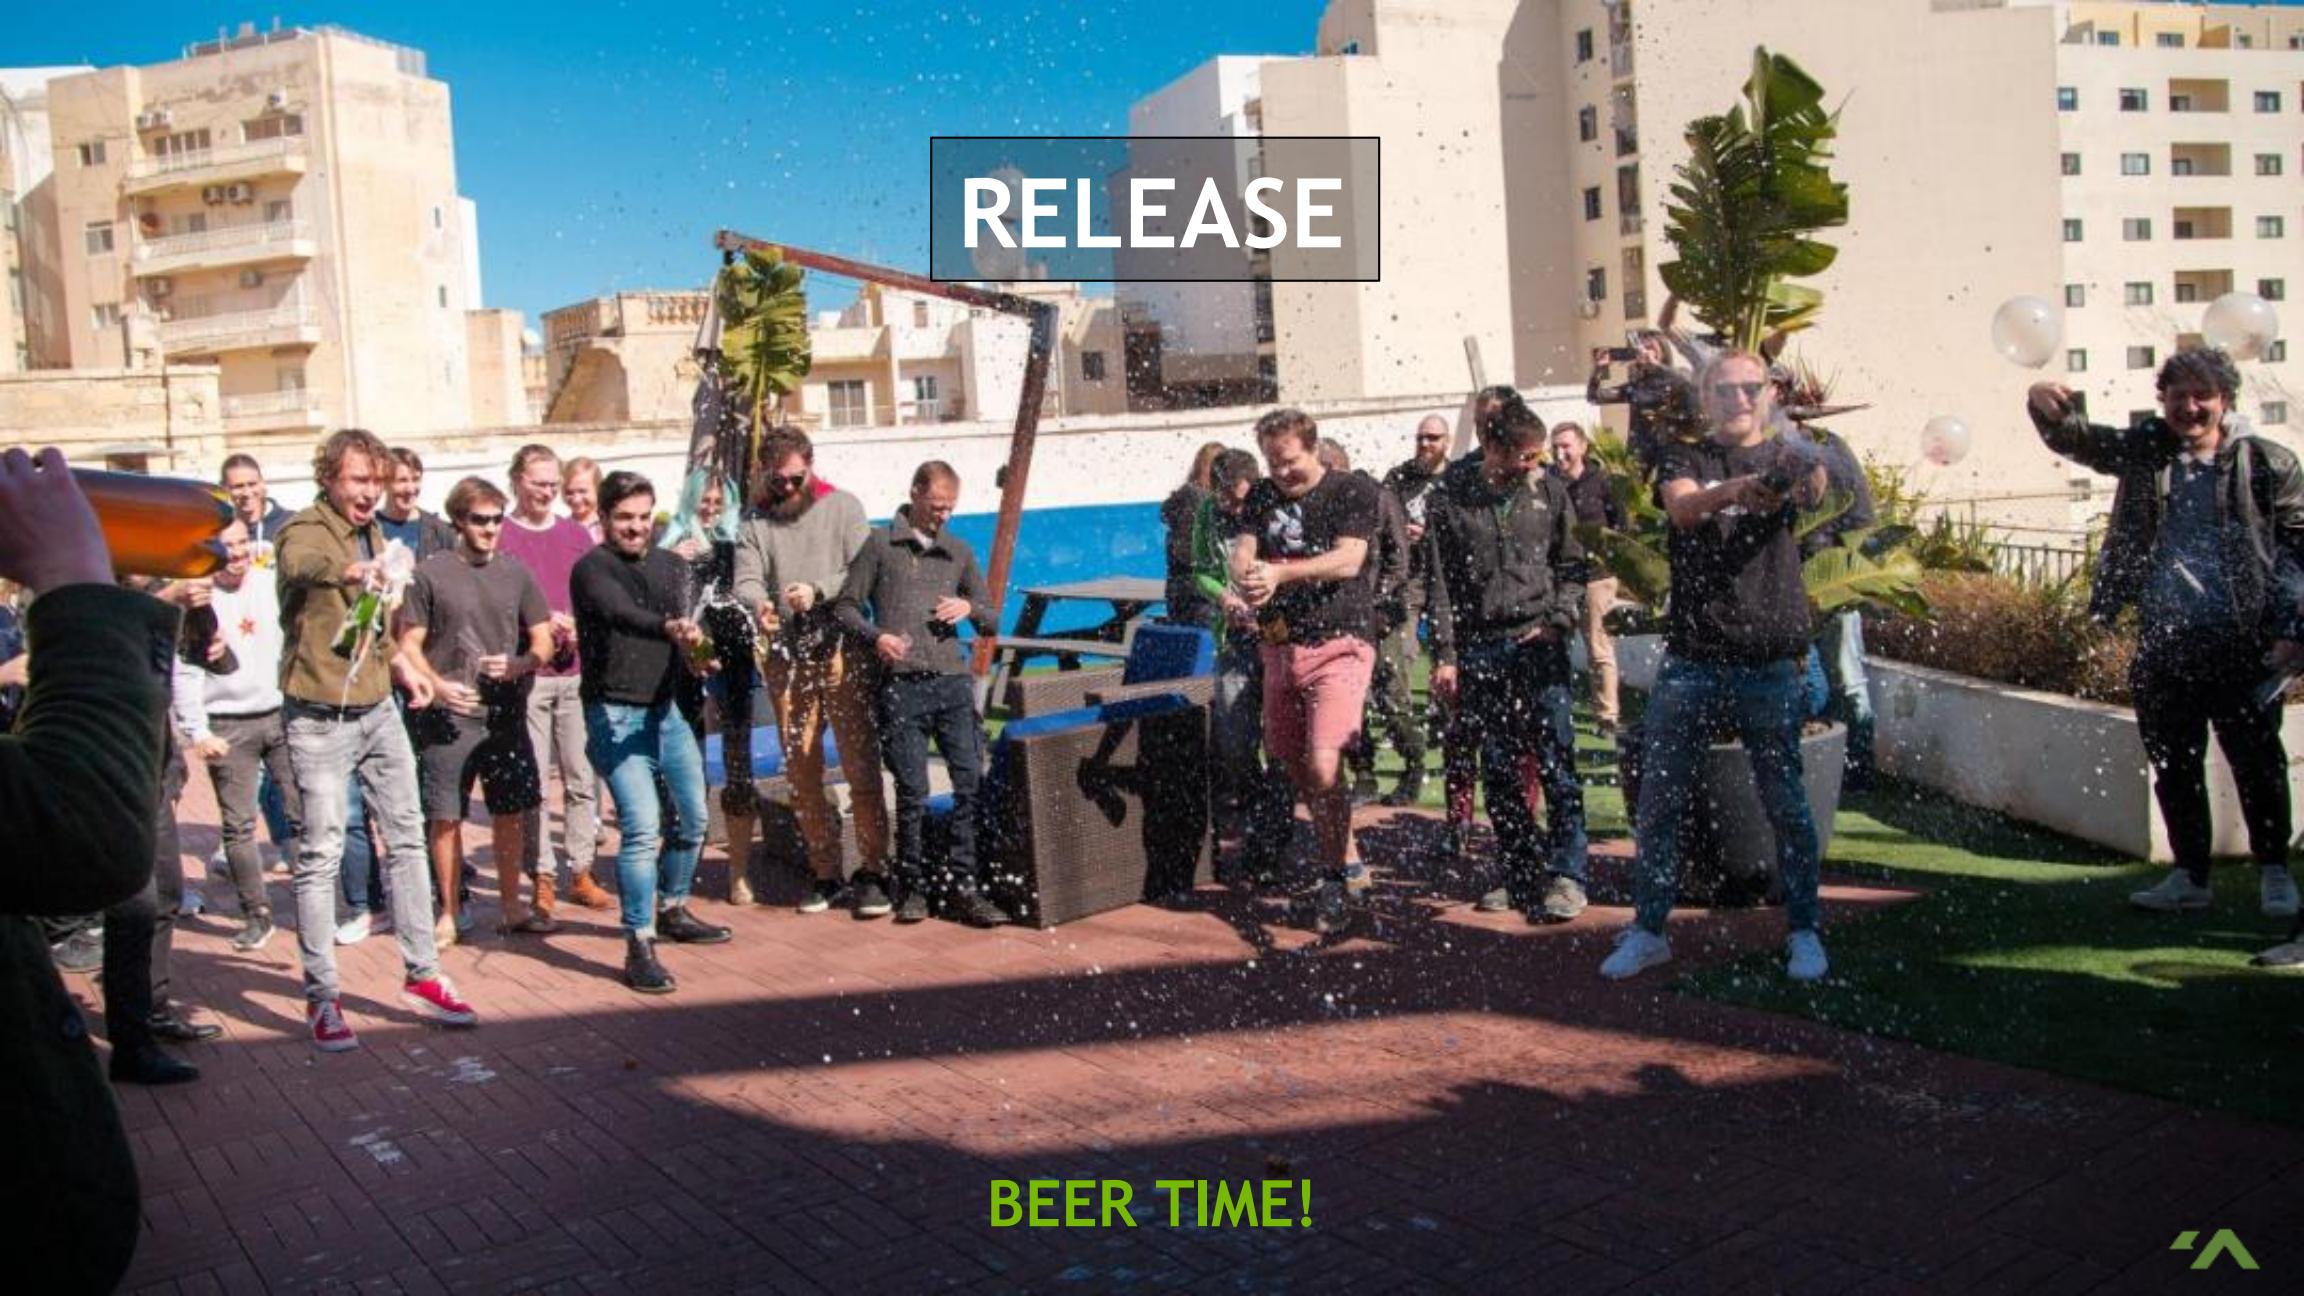
Adjust non-RTX mode content to match RTX in extreme cases

Both versions must look good!

There cannot be a loser  
it's Exodus vs Exodus



RELEASE



BEER TIME!



# NEW MEASUREMENT OF 'BETTER'

Enough of concerns

We do not expect RT-lighting to  
be exactly 'better'

Especially in an art-directed game

Results are clearly different

Mathematically stable solution  
makes them believable and  
natural

Or just convincing





RTX ?





RTX ?



# HOW RT MAKES US HAPPY

A tool to play with

An achievement

Fully dynamic solution - 4A's pillar

Lighting reference tool

Emergent results







# OUR NEXT DREAMS

What would Oles dream of next?

AO and GI are nailed

Area lights with soft shadows

Caustics. Magic in real life

Raytracing as one unified  
solution

Light-based gameplay logic

Deferred+Forward

Volumetrics

RT on consoles



# USEFUL LINKS

<https://media.contentapi.ea.com/content/dam/eacom/frostbite/files/gdc2018-precomputedglobalilluminationinfrostbite.pdf>

<http://orlandoaguilar.github.io/sh/spherical/harmonics/irradiance/map/2017/02/12/SphericalHarmonics.html>

[http://cg.ivd.kit.edu/publications/2017/svgf/svgf\\_preprint.pdf](http://cg.ivd.kit.edu/publications/2017/svgf/svgf_preprint.pdf)

[https://cg.ivd.kit.edu/publications/2018/adaptive\\_temporal\\_filtering/adaptive\\_temporal\\_filtering.pdf](https://cg.ivd.kit.edu/publications/2018/adaptive_temporal_filtering/adaptive_temporal_filtering.pdf)

# Спасибо!

Slides at [bit.ly/4agames](http://bit.ly/4agames)



Oles Shyshkovtsov | [oleksandr.shyshkovtsov@4a-games.com.mt](mailto:oleksandr.shyshkovtsov@4a-games.com.mt)

Sergei Karmalsky | [sergei.karmalsky@4a-games.com.mt](mailto:sergei.karmalsky@4a-games.com.mt)

Benjamin Archard | [benjamin.archard@4a-games.com.mt](mailto:benjamin.archard@4a-games.com.mt)

Dmitry Zhdan | [dzhdan@nvidia.com](mailto:dzhdan@nvidia.com)

# BONUS SLIDE!

## Color to spherical harmonics

```
float4 ConvertToIrradianceSH(float3 color, float3 dir, out float2 CoCg)
{
    float    Co      = color.r - color.b;
    float    t       = color.b + Co * 0.5;
    float    Cg      = color.g - t;
    float    Y       = max(t + Cg * 0.5, 0.0);

    CoCg = float2(Co, Cg);

    float    L00     = 0.282095;
    float    L1_1    = 0.488603 * dir.y;
    float    L10     = 0.488603 * dir.z;
    float    L11     = 0.488603 * dir.x;
    float4  shY     = float4 (L11, L1_1, L10, L00) * Y;

    return shY;
}
```

# BONUS SLIDE!

Spherical harmonics to color (no resolve)

```
float3 GetColorFromIrradianceSH(float4 shY, float2 CoCg)
{
    float Y      = shY.w / 0.282095;
    float T      = Y - CoCg.y * 0.5;
    float G      = CoCg.y + T;
    float B      = T - CoCg.x * 0.5;
    float R      = B + CoCg.x;

    return max(float3(R, G, B), 0.0);
}
```

# BONUS SLIDE!

## Spherical harmonics resolve

```
float3 ResolveIrradianceSHToDiffuse(float4 shY, float2 CoCg, float3 N)
{
    float d = dot(shY.xyz, N);
    float Y = 2.0 * (1.023326 * d + 0.886226 * shY.w);
    Y = max(Y, 0.0);

    // correct color-reproduction
    CoCg *= 0.282095 * Y / (shY.w + 1e-6);

    // YCoCg -> RGB
    float T = Y - CoCg.y * 0.5;
    float G = CoCg.y + T;
    float B = T - CoCg.x * 0.5;
    float R = B + CoCg.x;

    return max(float3(R, G, B), 0.0);
}
```